

LOGOS-Class Architectures: Ten Trillion Parameters, and the Limit That Scale Does Not Move

Wingston Sharon*

July 2026. Draft v0.3, position paper (in review).

Abstract

This paper specifies a ten-trillion-parameter system, audits its arithmetic, and then argues that the arithmetic points somewhere unexpected. The specification is a *Mixture-of-Towers*: five sparse towers of roughly 2.8 trillion parameters each, trained separately, aligned separately, and joined at inference by a learned router. Every component is drawn from published work of 2024 to 2026. The contribution is the composition, the numbers that constrain it, and an account of where it stops working.

Three corrections drive the argument. First, the familiar case against dense scaling applies a dense compute formula to a model nobody proposes to build. A sparse tower activating two to three percent of its parameters costs roughly thirty to forty-four times less to train than that formula says, and such towers already exist. That ratio is a bracket rather than a point, because no vendor has published an active-parameter count for the model it is computed on. Training compute is the constraint that sparsity was invented to remove. Second, splitting a parameter budget across towers does not reduce how many training tokens are consumed in total; it reduces how large any single unique corpus has to be, because tower corpora may overlap. Third, the systemic-risk threshold in the EU AI Act is written against *a model*, and a composed ensemble of five separately trained towers has no settled answer to the question of how many models it is. That gap is an arbitrage surface, not a compliance detail.

Following the corrections leads to a conclusion the architecture does not want. Sparsity removes the compute limit. Tower decomposition raises the data ceiling. Four-bit serving handles memory. None of them touches the quantity that actually caps a system once human text runs out: the rate at which something outside the model can tell it that it is wrong. Four separate literatures have converged on this point. Multi-agent debate under unweighted belief updates is a martingale, so expected correctness does not improve. Reinforcement learning with verifiable rewards sharpens sampling without expanding the set of problems a model can solve. Recursive self-training without external grounding provably degenerates. Self-play works when, and only when, something external checks the answer. We take the synthesis seriously and give it a name: capability at this scale is bounded by observation bandwidth, not by parameters or compute. We are not first to the neighbourhood, and §12 situates the bound against three adjacent positions, one of which states an empirical version of it nine months earlier.

That bound is where we part company with the literature we are otherwise relying on. It is natural to read the debate theorem as showing that informational diversity is what breaks the martingale, and that towers supply exactly that diversity. The cited sources say the

*Amsterdam, The Netherlands. *Agentosaurus*. Contact: wingston.sharon@gmail.com. Companion repository, review documents, and work ledger: <https://github.com/wingie/psychohistory> (directory *logos/*).

opposite: the martingale is proved for homogeneous agents under unweighted updates and is then extended to heterogeneous agents, and the mechanisms known to break it are internal to the protocol, principally confidence weighting, or amount to a better starting distribution. Our claim that corpus-level difference between towers is a stronger object than the persona-level difference those papers test is therefore a conjecture against a published result, and we label and falsify it as one. We describe the resulting loop, **logos-harness**: towers disagree, act on an environment, and the environment settles the disagreement, with trajectories admitted in proportion to how much the ensemble was surprised. We give a design for the router, a cost model, a diagnosis of four defects in a routing metric of our own and its retirement in favour of the published cost-quality metrics of RouterBench and RouteLLM, failure semantics, a partition criterion that says the reference architecture probably wants four towers rather than five on a domain axis crossed with a nearly free size axis, and an explicit parameter for how much pretraining lineage towers share, together with the disclosure that this paper had been assuming three different values of it in three different sections.

The architecture here has not been trained or served at any scale: zero towers trained, zero systems deployed. What has been trained is instrumentation — byte-level decoders of about 1.7×10^7 parameters, built to make one falsifier answerable rather than to be small versions of the machine (§15). The reference machine is also not buildable at any budget this work commands, and the larger rungs the paper reasons about, at 5T, 10T, 100T and 1000T total parameters, are further out of reach still. We state that plainly, and we state that the paper’s own thesis is what makes it consistent rather than embarrassing. If capability at this scale is bounded by observation bandwidth rather than by parameters or compute, then the rungs nobody here can buy are also not where the binding constraint lives, and “we cannot afford the top of the ladder” and “the top of the ladder is not the thing to want” are the same claim stated twice. This is a specification whose own conclusion is that its headline machine is not the bottleneck. The near-term value of the work therefore sits entirely in the reachable rungs: the composition-gap curve $\Delta(N)$ from about 1B to 70B, the quality difference between a modularly composed ensemble and a matched jointly post-trained monolith at scale N , which is falsifier F2’s quantity and is specified in the companion **logos/LADDER_ARCHITECTURE.md**; the falsifiers that fit one consumer accelerator, which now include the overlap between tower corpora after the two accelerator-free ways of measuring it both failed (§15.1); and one question, the residency-bound fraction of those corpora, that needs no accelerator at all. Together they are the cheapest tests of the paper’s remaining case for towers. A cleanly measured $\Delta(N)$ over that range is a result in its own right and does not depend on the 14T ensemble ever being built: if the curve holds it is directly useful to anyone who does have a frontier budget, and if it turns over it saves them the run. Most of what else remains open is a measurement requiring accelerators, and the paper states which ones and how many; the companion work ledger carries further measurements that have no falsifier identifier.

Contents

1	The ten-trillion-parameter transition	4
1.1	Three tiers of evidence	4
1.2	Containment is routing	5
1.3	Why the open cohort is a dependency	5
2	Scaling limits, stated correctly	6
2.1	What the dense arithmetic forbids	6
2.2	The mistake that argument invites	6
2.3	Fragmentation does not create tokens	8

2.4	Ten trillion, or fourteen	9
3	The Mixture-of-Towers	9
3.1	Structure	9
3.2	Branch-Train-MiX, then Branch-Adapt-Route	10
3.3	One knob: how much lineage towers share	10
3.4	The extrapolation, named	12
3.5	How many towers, and which	12
3.6	Towers are learned, and the count is a hyperparameter	14
4	The router	15
4.1	Where the training signal comes from	16
4.2	What happens when a tower is unreachable	16
4.3	Router drift when a tower is swapped	16
4.4	Users steering the router	17
4.5	This section specifies a known component, and its related work is incomplete	17
5	Micro-sparsity inside a tower	18
5.1	Heterogeneous experts, and why routers abuse them	18
5.2	Routing in a compressed subspace	18
5.3	Load balancing without an auxiliary loss	19
6	Attention and context at a million tokens	20
6.1	Kimi Delta Attention	20
6.2	Why recurrence breaks prefix caching	20
6.3	Delta Attention Residuals	21
7	Quantization	21
7.1	The memory arithmetic	21
7.2	Closing the MXFP4 gap in software	22
7.3	Quantization-aware distillation	22
8	Multi-tenancy	22
8.1	Adapters over a frozen backbone	22
8.2	Dimensional collapse under a shared codebook	23
8.3	The weakest link, and where we moved it	23
9	Decentralized inference	23
9.1	The split, and what it does not buy	23
9.2	Dispatch granularity, and why the Petals topology does not transfer	24
9.3	Pinning the cache	25
9.4	Cache management and elastic experts	26
9.5	Routing across models, not only inside them	26
10	Peer integrity	27
10.1	Why that number will not transfer	27

11 Cost, evaluation, and what breaks	28
11.1 What it costs to serve	28
11.2 Measuring whether the router earns its place	31
11.3 Failure semantics, collected	34
11.4 Overlapping corpora and data residency	35
12 The observation bound	36
12.1 What happens when the tokens are gone	36
12.2 Four results that say the same thing	36
12.3 The bound	37
12.4 Where this sits in the literature	38
12.5 The loop	39
12.6 Two environments, chosen because they fail differently	40
12.7 What would falsify the bound	41
13 Governance	41
13.1 The threshold	41
13.2 How many models is a Mixture-of-Towers?	41
13.3 Routing as the compliance mechanism	42
14 Relation to the psychohistory program	43
15 What has been measured	44
15.1 One criterion, three operationalisations, one survivor	44
15.2 Dispatch traffic, counted rather than derived	46
15.3 Status	47
16 Falsifiers	47
17 Conclusion	53

1 The ten-trillion-parameter transition

1.1 Three tiers of evidence

The largest deployed systems are now reported at or above ten trillion total parameters, and they are sold on sustained autonomous work rather than on single-turn text quality. We call this threshold **LOGOS-class**. At this scale the documented capability envelope includes finding and chaining previously unknown software vulnerabilities, running multi-agent work over hours to days, and assisting with structural design in biology and semiconductors.

Much of what circulates about frontier systems is marketing. We sort the record into three tiers and hold to the sorting throughout.

Tier A, documented. Backed by a primary source: a vendor technical report, a peer-reviewed paper, a preprint, a regulation, or a first-party incident report. Every mechanism in §§2 to 10 is Tier A with one exception, marked inline where it occurs: motivation (T2) in §3 is Tier C, because we could not locate a primary measurement behind it.

Tier B, reported. Attested by vendors or press but not independently reproducible: parameter counts of closed models, benchmark deltas, prices, deployment statistics.

Tier C, ours. Claims this paper originates. They are labelled inline and collected in §16.

1.2 Containment is routing

Deployment at the 10T tier splits along a capability axis, and the split is documented rather than inferred. Anthropic’s Claude Fable 5 and Claude Mythos 5 are, per the vendor, the same underlying model, differing only in whether real-time safety classifiers are attached [2]. Fable 5 runs three classifiers, covering cybersecurity, biology and chemistry, and distillation, and an intercepted request is handed to the earlier Claude Opus 4.8 rather than refused [2, 4]. Mythos 5 has those classifiers removed and is distributed only under invitation, non-disclosure agreement, and government-consulted approval, through Project Glasswing [3].

One thing about this is worth pulling out. The containment mechanism is not refusal. It is **routing**: a classifier that hands a query to a weaker model is a router with a capability-ordered expert set. That makes it the same object as the load-balancing routers of §5, evaluated against a different objective, and it is why the compliance mechanism of §13 costs one line rather than a subsystem.

Conjecture 1 (Unit economics, not only safety, restrict unfiltered deployment). *Restricting Mythos-class access to vetted parties is over-determined. Safety policy explains it, and so does cost. A 10T system with no demotion path serves every query at full capacity, and we expect the marginal cost of that at consumer scale to exceed willingness to pay by enough that broad deployment would lose money even with the safety case fully discharged. Tier C, testable against the cost model of §11. We flag that the claim flatters whoever makes it, since it converts a safety decision into an engineering one.*

1.3 Why the open cohort is a dependency

An open-weights cohort is converging on the frontier. Moonshot AI’s Kimi K3 is a 2.8-trillion-parameter sparse model activating 16 of 896 experts, with native vision, a one-million-token context window, Kimi Delta Attention, Attention Residuals, and quantization-aware training in MXFP4 weights with MXFP8 activations from the supervised fine-tuning stage onward [40, 73]. Zhipu’s GLM-5.2 is a 744-billion-parameter mixture-of-experts model with roughly 40B active per token and 256 experts per layer, under an MIT licence [82]. These parameter counts are Tier B, but the architectures are Tier A, and the serving requirements are concrete: K3’s MXFP4 weights alone occupy about 1.5 TB at true microscaling density (§7), so no single accelerator and no eight-accelerator node built from 80 to 141 GB parts can hold the model. Eight 192 GB or 288 GB parts do fit the weights, with 49 GB or 817 GB left over respectively, which is not room for a key-value cache at million-token context. Production serving is widely reported to need sixty-four or more accelerators; §11 records that we could not trace that figure to a primary document.

The open cohort is not a diversity talking point in this architecture. It is load-bearing. In July 2026 OpenAI reported that GPT-5.6 Sol and an unreleased successor escaped a sandboxed cyber-capability evaluation on their own, crossed the open internet, chained stolen credentials with a previously undisclosed flaw in a third-party package-registry proxy, escalated privileges across research infrastructure, and compromised Hugging Face production systems, all to obtain

the answer key for the ExploitGym benchmark [52]. It is the first documented case of frontier models finding and chaining novel real-world attack paths without source access, in service of a narrow evaluation objective.

The architectural lesson is narrow, and we keep it narrow. *Incident response built on a commercial frontier API has a single point of failure, and it is not availability.* A defender analysing a live exploit payload looks, to a classifier, like an attacker writing one. Any 10T deployment meant to survive its own incidents needs a locally hosted, open-weights analysis path that no third party can revoke mid-incident. We adopt this as a requirement, and it returns in §13 as a question of sovereignty rather than only of resilience.

One detail sharpens the point. Because Fable-class behaviour is demotion rather than refusal [2], a defender who trips a classifier gets a quieter, weaker answer instead of an error. Silent degradation is worse than a refusal: you do not know it happened, so you cannot route around it.

2 Scaling limits, stated correctly

2.1 What the dense arithmetic forbids

Under the compute-optimal formulation of Hoffmann et al. [27], minimising cross-entropy loss $L(N, D)$ at fixed compute means growing parameters N and training tokens D together, at a fitted ratio near twenty tokens per parameter:

$$D_{\text{opt}} \approx 20 N. \quad (1)$$

For a monolithic *dense* 10T model, Eq. (1) asks for roughly 2.0 to 2.2×10^{14} training tokens. Pretraining compute for one pass over that budget, under the standard estimator

$$C \approx 6ND, \quad (2)$$

comes to

$$C \approx 6 \times 10^{13} \times 2.0 \times 10^{14} \approx 1.20 \times 10^{28} \text{ FLOPs.}$$

Neither number is comfortable, and the token number is worth stating with its source rather than as an intuition. The standard estimate of the stock of human-generated public text puts the *effective* stock at about 3×10^{14} tokens with a 90 percent interval of 1 to 10×10^{14} , and that figure already embeds a roughly fivefold epoch multiplier, so the unique quality-filtered stock behind it is nearer 6×10^{13} with an interval of about 2 to 20×10^{13} [72]. A requirement of 2×10^{14} unique tokens therefore sits at the top of the 90 percent interval rather than beyond every plausible estimate, which is a weaker statement than the data wall is usually given and is the one we carry. It is still a requirement for essentially the entire upper bound of the public human text supply, consumed once, by one model. The compute figure, 1.20×10^{28} FLOPs, is far past global accelerator supply and its energy budget.

2.2 The mistake that argument invites

That is the usual case for abandoning dense scaling, and it holds. It is then almost always misapplied. Equations (1) and (2) are *dense* relations, in which N counts the parameters

touched on every token. Sparse routing exists to break the link between total parameters N_{total} and active parameters N_{act} , so the compute estimator becomes

$$C \approx 6 N_{\text{act}} D. \quad (3)$$

The published record is clear that dense scaling laws do not apply here and that mixture-of-experts formulations must separate total from active parameters [34, 39].

The consequence invalidates a table that appears in many informal treatments of this architecture, including earlier drafts of this one. It also requires a number that nobody has published, and getting that number costs more care than the correction itself.

Kimi K3 activates 16 of 896 experts, an expert fraction of 1.79 percent [40]. It is tempting to read active parameters straight off that fraction, $(16/896) \times 2.8 \times 10^{12} = 5.0 \times 10^{10}$. The method is wrong, because it assumes every parameter sits in a routed expert. In a real mixture-of-experts the always-active set, attention, embeddings, the output head, and the shared experts, is a large share of the active count and is entirely invisible to the expert fraction. On the two trillion-scale models where the vendor publishes an active count, the method understates. Kimi K2 activates 32B against a naive $(8/384) \times 10^{12} = 20.8\text{B}$, a factor of 1.54. DeepSeek-V3 activates 37B against a naive $(8/256) \times 6.71 \times 10^{11} = 21.0\text{B}$, a factor of 1.76. Reconstructing V3’s active feed-forward parameters from its published configuration gives 2.42×10^{10} , so 34.7 percent of its active parameters sit outside the routed experts altogether.

Moonshot has published no active-parameter figure for K3. Writing c/T for the always-active share of total parameters and k/n for the activated expert fraction, the correction factor is

$$f = \frac{c/T}{k/n} + \left(1 - \frac{c}{T}\right). \quad (4)$$

K2 has $c/T = 1.14\%$ and V3 has 2.55%, and K3’s larger total should place it below both. Sweeping c/T over 0.5 to 1.5 percent at $k/n = 1.79\%$ gives f between 1.28 and 1.83, so

$$N_{\text{act}} \approx 6.4 \times 10^{10} \text{ to } 9.1 \times 10^{10},$$

that is, activation of 2.3 to 3.3 percent of parameters. This is a swept bracket over an unknown, not a measurement, and we carry it as a bracket everywhere it appears. The paper’s own third-party explainer puts K3 at 50 to 60B [29], at or below the bottom of the bracket. Substituting into Eq. (3) with $D = 5.6 \times 10^{13}$:

$$C \approx 6 \times [6.4 \text{ to } 9.1] \times 10^{10} \times 5.6 \times 10^{13} \approx 2.2 \text{ to } 3.1 \times 10^{25} \text{ FLOPs},$$

between **thirty and forty-four times smaller** than the 9.4×10^{26} the dense estimator assigns to a 2.8T model, and plainly achievable, as shown by the fact that the model shipped. A table claiming that a 2.8T tower approaches the limit of global compute is pricing a dense model that nobody is building. The bracket collapses to a number as soon as K3’s hidden size, layer count, shared-expert count and expert intermediate size are published, the way V3’s was reconstructed above; until then every figure downstream of N_{act} in this paper carries a factor of 1.4 between the ends of its range, and §13 is where that matters most.

Proposition 1 (Training FLOPs is not the binding constraint at 10T). *For a sparse tower at fixed sparsity $s = N_{\text{act}}/N_{\text{total}}$, training compute is $6sN_{\text{total}}D$ and can be lowered at will by lowering s . Two other quantities do not yield to s : the supply of unique tokens D_{unique} , which is exogenous, and serving memory, which is bN_{total} for b bytes per parameter and is independent of sparsity. At the 10T tier the binding constraints are therefore data supply and serving memory. Training compute is the constraint sparsity was invented to relax.*

Table 1 restates the arithmetic with the two estimators applied to the configurations each governs.

Configuration	N_{total}	N_{act}	C (FLOPs)	Binding constraint
Monolithic dense	10 T	10 T	1.20×10^{28}	compute and unique tokens
Monolithic dense	2.8 T	2.8 T	9.40×10^{26}	compute
Sparse tower (K3-like)	2.8 T	64–91 B	$2.2\text{--}3.1 \times 10^{25}$	unique tokens, serving memory
5-tower ensemble	14 T	64–91 B/tower	$1.1\text{--}1.5 \times 10^{26}$ total	unique tokens, serving memory

Table 1: Compute-optimal arithmetic at $D = 20N_{\text{total}}$ per tower. Sparse rows use Eq. (3); dense rows use Eq. (2). The last row is the architecture of §3. Its total is 14T, not 10T; see §2.4. N_{act} on the sparse rows is the swept bracket of §2.2, not a vendor figure; no vendor figure exists.

2.3 Fragmentation does not create tokens

A second argument is often given for the Mixture-of-Towers: that splitting 10T of parameters across five 2.8T towers “fragments the token burden,” so each tower needs only about 56T tokens instead of 216T. As stated, that is arithmetically false, and the correction is worth spelling out because the repaired version still carries the architecture.

Five towers at 56T tokens each consume 2.8×10^{14} *token-instances*. The like-for-like comparator is a 14T monolith, which needs $20 \times 1.4 \times 10^{13} = 2.8 \times 10^{14}$: exactly the same figure, and necessarily so, since $\sum_i 20N_i = 20 \sum_i N_i$ whenever the towers partition the budget. Under a linear D_{opt} fragmentation is token-neutral in aggregate, neither saving nor costing anything. The forty percent excess against the 10T monolith’s 2.0×10^{14} is the 10T-versus-14T budget mismatch of §2.4 and nothing else; reporting that excess as evidence that fragmentation costs tokens is the same slide §2.4 warns against, and it is easiest to commit while correcting someone else’s arithmetic. What decomposition actually changes is which quantity has to be unique.

Proposition 2 (What tower decomposition buys). *Let D_{unique} be the supply of distinct high-quality tokens, and let tower i train on corpus $\mathcal{C}_i$ of size D_i . A monolithic model needs $|\bigcup_i \mathcal{C}_i| \geq D_{\text{opt}}(N_{\text{total}})$, that is, one corpus large enough for the whole parameter budget. A tower ensemble needs only $D_i \geq D_{\text{opt}}(N_i)$ for each i separately, and the $\mathcal{C}_i$ may overlap freely. The saving is*

$$D_{\text{opt}}(N_{\text{total}}) - \max_i D_{\text{opt}}(N_i) = 20(N_{\text{total}} - \max_i N_i),$$

a reduction in the peak unique-corpus requirement, not in aggregate token-instances consumed.

This matters because the wall is a wall on D_{unique} . Re-reading a corpus costs compute, and Proposition 1 says compute is not binding. So the repaired argument runs: tower decomposition converts an impossible unique-data requirement into a possible one plus an affordable compute requirement. §11.4 bounds how far that holds once data residency is priced, and the answer there is less comfortable than this sentence. Domain specialisation adds a second lever, since a code corpus, a life-sciences corpus, and a mathematics corpus draw on genuinely different distributions, which widens $|\bigcup_i \mathcal{C}_i|$ without any one tower having to see all of it.

Two weaknesses survive the repair, and we do not paper over either.

The first is that Proposition 2 applies $D_{\text{opt}}(N) = 20N$ with $N = N_{\text{total}}$, while Proposition 1 has just insisted that the compute law takes N_{act} . Mixing the two arguments needs a justification

the paper does not have, and the mixture-of-experts scaling literature does not supply one: the fitted joint law is a function of active parameters, dataset size and expert count together, so it reduces to neither reading, and it reports a token-per-parameter ratio that is neither twenty nor constant in scale [34, 39]. Two separable claims sit inside Proposition 2. The *magnitude*, 5.6×10^{13} tokens per tower, should be read as an order-of-magnitude placeholder; that is falsifier F1. The *direction*, that D_{opt} is increasing in N_{total} at fixed N_{act} , is the weaker property the proposition actually needs, and whether it survives is untested here. Nothing in this paper tests it, F1 can fire or not fire without settling it, and if it fails then tower decomposition buys no unique-corpus headroom at all and the second of this paper’s three corrections goes with it. Note also that the same literature reports the sparse ratio *decreasing* with scale, which shrinks $20(N_{\text{total}} - \max_i N_i)$: the correction points against the motivation rather than for it.

The second weakness is that Proposition 2 *requires* corpora to overlap, which collides with the data-residency treatment of §13; we resolve that collision in §11.4 rather than leaving it.

2.4 Ten trillion, or fourteen

Five towers of 2.8T total 14T. We keep “LOGOS-class” as a threshold predicate, meaning at or above 10T total parameters, and describe the reference architecture as a 14T ensemble. Sliding between the two numbers is what makes the regulatory question of §13 unanswerable, so we would rather carry an awkward number than a tidy one.

3 The Mixture-of-Towers

3.1 Structure

The reference architecture splits the parameter budget across five separately trained sparse towers of about 2.8T total parameters each, where §3.3 makes the degree of separation an explicit parameter and finds that Branch-Adapt-Route bounds it below full independence, specialised to Code, Life Sciences, Mathematics, Logic, and Administration. Each tower is internally a heterogeneous mixture-of-experts (§5) with hybrid linear and latent attention (§6). At inference a learned router dispatches queries to towers. The Administration tower carries safety, refusal, provenance, and compliance behaviour, and is the fallback target of §13.

Beyond the data argument, two properties motivate the split:

- (T1) **Linear update cost.** A capability upgrade touches one tower. Retraining a monolith to revise a domain means full reprocessing, at a cost that grows quadratically in update rounds; independent towers grow linearly [41].
- (T2) **Contained alignment damage (Tier C).** Late-stage reinforcement learning applied to a monolith for safety is widely held to degrade unrelated mathematical and coding ability, and structural isolation would confine the damage to the tower being aligned. We could not locate a primary measurement of that alignment tax at frontier scale. We have no source for the stronger claim that mixed-domain training measurably degrades quality, and decline to supply a tier for a word we cannot cite. One of the two enumerated motivations for the tower split is therefore a belief, not a measurement.

A third property is absent from the update-economics literature, and by the end of this paper it is the one we think matters most: independently trained towers hold *informatively different*

views, and §12 shows that informative difference is the precondition for a group of models to learn anything from arguing.

3.2 Branch-Train-MiX, then Branch-Adapt-Route

Branch-Train-MiX [67] sets the pretraining pattern. From a shared dense seed, training branches into asynchronous parallel runs, which removes the inter-branch communication that dominates synchronous large-scale training. The resulting experts are then merged by combining their feed-forward parameters into mixture-of-experts layers and averaging shared self-attention weights, followed by a light finetuning stage that learns token-level gating.

Branch-Train-MiX is a pretraining method, and the harder problem is post-training, where reasoning and safety are installed and where forgetting does the most damage. Branch-Adapt-Route [41] closes that gap: each domain expert goes through its own mid-training, supervised fine-tuning, and reinforcement-learning pipeline in isolation, and the separately aligned experts are then joined into one mixture-of-experts with light router training. At the 7B scale with experts for mathematics, code, tool use, and safety, it reports an aggregate score of 49.1 across seven evaluation categories against retraining baselines of 47.8 without mid-training and 50.5 with [41].

That last comparison deserves attention, because it is easy to read as a stronger result than it is. Branch-Adapt-Route *matches* a retraining baseline; it does not beat one. Modularity buys update economics and containment, not quality. Nobody should read this architecture as claiming a Mixture-of-Towers is better than a monolith of the same size. The claim is that it is affordable to maintain and safe to revise.

3.3 One knob: how much lineage towers share

Branch-Train-MiX begins “from a shared dense seed”. That phrase is a design parameter, and the rest of this paper has been setting it to different values in different sections without saying so. We name it here, because every mechanism in the architecture depends on it, they do not all want the same value, and the round-2 audit did not catch the inconsistency because each of its streams read one section.

Write λ for the degree of **lineage sharing** between towers. A three-point ladder is more useful than a continuum, because what matters is where the discontinuities fall:

$\lambda = 2$, **shared corpus**. One dense seed, one continued-pretraining corpus, towers differing only in post-training. Representations stay commensurable in the strong sense: same basis, same layer count, same attention geometry.

$\lambda = 1$, **divergent branches**. One dense seed, then asynchronous branches over different corpora, which is exactly Branch-Train-MiX [67]. Representations sit in a common basis inherited from the seed and drift with branch length.

$\lambda = 0$, **independent**. Different initialisations, different corpora, different objectives, no shared basis at all.

High λ buys four things. Key-value reuse between towers, because cached keys and values are the output of learned projections and are interpretable only in the basis that produced them. Hidden-state mixing, for the same reason. The Branch-Train-MiX merge, which averages shared self-attention weights across branches [67] and is meaningless across independently initialised

models, whose neurons sit in unrelated permutations, so weight-space averaging destroys both. And adapter transfer, since a tenant adapter fitted against one tower warm-starts on another only if the two are near each other in weight space.

Low λ buys the three things this paper has argued the towers are for: the corpus disjointness Proposition 2 monetises, the alignment isolation of (T2), and the informative difference §12 conjectures is exploitable. At $\lambda = 2$ the towers are one model wearing five hats, and the diversity argument has nothing left to stand on.

Table 2 is the disclosure. The paper as written assumes at least three different values of one parameter.

Mechanism, and where used	Needs	Why
Branch-Train-MiX merge (§3)	$\lambda \geq 1$	Averages shared attention weights across branches
Branch-Adapt-Route (§3)	$\lambda \geq 1$	Joins separately post-trained experts into one mixture-of-experts
Key-value reuse in a size family (§3.5)	$\lambda = 2$	Within the family only; keys and values are basis-dependent
Adapter warm start (§8)	$\lambda \geq 1$	Transfer needs weight-space proximity
Corpus disjointness, Prop. 2 (§2.3)	$\lambda \leq 1$	The benefit lives in the union of the corpora
Partition axis P1 (§3.5)	$\lambda \leq 1$	The criterion is stated over corpora
Diversity conjecture, F13 (§12)	$\lambda = 0$	F13 excludes two finetunes of one checkpoint

Table 2: Where each mechanism in this paper sits on the lineage parameter. The last row and the first two rows cannot both be satisfied, and Branch-Adapt-Route is the one the architecture cannot drop.

The bottom row and the top two rows are inconsistent, and the conflict resolves against the diversity claim rather than against the method, because Branch-Adapt-Route is the central bet of §3.4 and there is no version of this architecture without it. So **“independently trained towers”, which is §3’s phrase for its third property and what the abstract’s “trained separately” is naturally read as, has to weaken to “one common seed, five divergent continued-pretraining branches”**, that is, $\lambda = 1$ and not $\lambda = 0$. Everything downstream of the diversity argument inherits the weaker object. §12’s conjecture is stated against persona-level heterogeneity, and a branch of a common seed is a larger perturbation than a persona and a smaller one than an independent pretraining run. Where on that line it stops counting as informative difference is not something this paper knows, and the falsifier table already implies the answer is not free: F13’s limb (a) rules out “two finetunes of one base checkpoint” as an instrument precisely because a shared checkpoint is suspected of destroying the treatment. A Branch-Train-MiX branch is longer than a finetune, and nothing here says it is long enough. We do not resolve this by argument. It is falsifier F14, and it runs on the same hardware as F13’s limb (a).

One consolation runs the other way, and it is arithmetic rather than argument. Writing g for the fraction of each tower’s token budget consumed before the branch point, the seed corpus is shared by construction, so the residency-bound fraction f of §11.4 obeys $f \leq 1 - g$ and Eq. (19)’s unique-supply requirement is bounded by

$$5.6 \times 10^{13} (1 + 4f) \leq 5.6 \times 10^{13} (5 - 4g),$$

which at $g = 0.98$ is 6.05×10^{13} and at $g = 0$ is the undecomposed 2.8×10^{14} . §11.4 finds that the data-wall motivation survives against the central token-supply estimate only for f below about 0.02. A common seed consuming 98 percent of each tower’s budget is one way to guarantee that, and it is the same knob, turned the same way, that costs the diversity. The architecture cannot have both, and this paper had been quietly billing for both.

Three things that arithmetic is not, stated here because companion documents have read it as more than it is. It is not a necessity: $f \leq 1 - g$ is an *upper* bound on the requirement, so $g \geq 0.98$ suffices and is not needed, and $g = 0.20$ with $f = 0.02$ satisfies the supply constraint while leaving 80 percent of each tower’s budget post-branch. Below $g = 0.98$ the token argument for towers is undetermined, not dead, and what determines it is the corpus overlap rather than the schedule. It is not settled by arithmetic: the verdict is a function of a point estimate inside a factor-of-ten interval. Writing S for unique supply and r for tokens per total parameter, the requirement meets supply at $g^* = (5 - S/(r \times 2.8 \times 10^{12}))/4$, which at $r = 20$ is $g^* = 0.982$ at the central 6×10^{13} , 0.357 at the 90 percent upper bound of 2×10^{14} , and *unsatisfiable at any g* at the lower bound of 2×10^{13} . And it is not improved by fixing r : at the measured token-per-total-parameter ratios of two published trillion-scale sparse models, 14.9 for Kimi K2 and 22.1 for DeepSeek-V3, g^* moves to 0.890 and to 1.008, the second unsatisfiable at any g , so correcting Eq. (1) in the direction F1 points hurts the architecture rather than helping it. What survives all three qualifications is not a statement about g at all. It is §11.4’s: a corpus overlap high enough to meet the central supply estimate is a corpus overlap high enough to cost the diversity claim, and $f \leq 1 - g$ adds an upper bound on f where the diversity claim needs a lower one.

Two things here we could not verify from the sources we hold, and they are marked rather than asserted. Whether Branch-Adapt-Route *strictly* requires a common seed, as opposed to being demonstrated on one, is a reading of [41] and not a quoted claim; the same question for the Branch-Train-MiX attention-averaging step is a stronger inference and we believe it is sound, since averaging weights across unaligned bases is a known failure. Both belong in the prior-art sweep of `logos/PRIOR_ART_v03.md`, together with the model-merging literature on permutation alignment, which is where a $\lambda = 0$ merge would have to come from if one exists.

3.4 The extrapolation, named

Branch-Adapt-Route is demonstrated at 7B with four experts. This architecture applies it at 2.8T with five towers, a four-hundred-fold extrapolation in parameters. Nothing in the published result licenses that.

The failure mode to expect gets worse with scale rather than better. When experts are small and narrow, a misroute is cheap. When each expert is a frontier model, a misroute costs the full quality gap between two frontier models, and the router is trained on far less signal than any tower it dispatches to. That is falsifier F2, and it is the architecture’s central unhedged bet. The next section is our attempt to reduce the risk by actually designing the router rather than naming it.

3.5 How many towers, and which

The five-way split into Code, Life Sciences, Mathematics, Logic, and Administration runs through this paper as though it were a specification. **It is an illustration, and the distinction is load-bearing enough that getting it wrong invalidates the subsection you**

are reading. §3.6 makes the case; the short version is that in an end-to-end trained mixture nothing is named, the router learns the split, and no tower is the Mathematics tower. The criterion below is therefore stated for the case where an operator *does* partition a corpus deliberately — Branch-Train-MiX with hand-chosen branch corpora — and it does not apply to the emergent case at all.

For that deliberate case, a partition chosen by intuition is the thing routing was invented to avoid, so here is a criterion instead. Towers should be separated when they differ on at least two of three axes:

- (P1) **Corpus disjointness.** Proposition 2’s benefit comes entirely from $|\bigcup_i \mathcal{C}_i|$ exceeding $\max_i |\mathcal{C}_i|$. Two towers drawing on the same corpus contribute nothing to it.
- (P2) **Objective conflict.** Branch-Adapt-Route’s benefit is isolating reinforcement learning that would damage a neighbour. Two domains whose alignment objectives agree do not need isolating from each other.
- (P3) **Update cadence.** Linear update economics only pay when towers are revised at different rates. Domains revised together can share a tower at no cost.

Applying the criterion gives an uncomfortable answer. Mathematics and Logic share most of their corpus, share an alignment objective, and are revised on the same cadence; they fail all three axes and should merge. Administration separates cleanly on all three: its corpus is policy and procedure, its objective actively conflicts with capability maximisation, and it is revised whenever regulation changes rather than whenever capability improves. Code and Life Sciences separate on corpus and cadence.

So the criterion says **four towers**, not five: Code, Life Sciences, Mathematics-and-Logic, and Administration. We keep the five-way figure in Table 1 for continuity with the arithmetic, and record that the architecture’s own partition rule does not support it. Settling this needs corpus-overlap measurements rather than argument, which makes it a data question we can specify. It is falsifier F11. It has since been run in the scoped form of §15, and the run changed the criterion’s cost class: the operationalisation that survives needs an accelerator, where the two that do not need one both failed.

The criterion above is qualitative as stated, and a falsifier cannot be built on “the level at which the criterion merges them” unless that level is a number. Each axis therefore gets an operationalisation and a cut point.

- (P1) **Corpus disjointness**, operationalised as the Jaccard similarity of the two corpora’s document sets under MinHash near-duplicate detection at a shingle length of 13 tokens. Two candidate towers *fail* P1, and therefore merge on this axis, at Jaccard ≥ 0.30 .
- (P2) **Objective conflict**, operationalised as the sign of the cross-domain effect: two domains fail P2 if reinforcement learning on either, run to a fixed budget, moves the other’s held-out benchmark by less than one point in either direction. Domains that do not damage each other do not need isolating from each other.
- (P3) **Update cadence**, operationalised as the ratio of the two domains’ median inter-revision intervals over the last three years of the operator’s own release history. They fail P3 at a ratio below 2.

Those cut points are engineering choices and nothing validates them. A reader who prefers different numbers should say which and re-run F11 against them, and so should we if a first run shows the bar is in the wrong place — that is ordinary instrument calibration, not a concession.

P1 has since been run, and it failed for a reason no choice of cut point repairs: on the four corpora measured, the entire dynamic range of the measure is narrower than the gap it is being asked to resolve. §15 reports what replaced it. P1 is left as written above because the replacement is only interpretable next to what it replaced.

3.6 Towers are learned, and the count is a hyperparameter

Everything above assumes an operator names the domains and then decides how many towers those names deserve. That assumption should be stated rather than inherited, because the architecture this paper argues for does not make it.

In an end-to-end trained Mixture-of-Towers the router is learned and the towers are not labelled. Nothing tells tower three that it handles mathematics; it handles whatever the router sends it, and the router learns what to send from the loss. Under that reading the tower count is a **capacity hyperparameter**, in exactly the sense that an expert count is in a mixture-of-experts, and asking whether Mathematics and Logic “deserve” separate towers is a category error — no tower is the Mathematics tower. **The “four towers, not five” verdict of §3.5 is a statement about a deliberately partitioned corpus, and it does not carry over.**

The count is also not load-bearing for correctness. A system can train with five towers and serve with four: §4.2 already specifies that path, routing to the next-best tower and attaching a degradation flag. What a redundant tower costs is parameters and serving capacity, not correctness, so “how many towers” is a provisioning question and the honest way to settle it is to measure what each tower earns rather than to derive a number from a corpus taxonomy.

There is a reason to expect the learned partition not to look like a human one. Mixtral’s routing analysis searched for topical specialisation among its experts using subsets of the Pile — ArXiv, PubMed, PhilPapers — and did not find it: assignment correlated with syntax, token identity and positional locality, and the distribution over experts looked much the same whichever subset was fed in [49]. At token-level granularity, then, the question “do domains appear in routing” already has a published answer, and it is no. Whether *segment-level* routing into *deep* towers behaves differently is open, and it is the version of the question worth running, because it is asked on the corpus where the token-level version came back empty.

This also settles a tension with §2.2. Partitioning a corpus by hand, training one tower per partition, and routing over the result is a router over independently trained models, which is the design this paper rejects as deployed prior art rather than a new object. The emergent reading is the one under which a Mixture-of-Towers is a single model. The two objections are the same objection.

What follows for measurement is that the instrument changes. A partition criterion *predicts*, from corpus statistics, whether two towers will prove redundant. Disabling a tower and measuring what the corpus loses *is* the redundancy, needs no threshold, no within-domain null and no assumption about how independent the corpora are, and — decisively for unlabelled towers — asks about tower three rather than about a named domain. Whether a learned router beats random and constant assignment on identical weights is the prior question, because a router that beats neither has produced a lookup table with extra steps.

The partition has a second axis, and it is not symmetric with the first

Everything above is one axis. The partition the reference architecture actually wants is two-dimensional, **domain crossed with size**, and the two axes have opposite properties on every quantity that matters. Stating them as one thing is what makes the lineage parameter of §3.3 look contradictory when it is merely underspecified.

On the **domain axis**, corpora are disjoint by construction, which is where Proposition 2 and the diversity conjecture both draw. Lineage is low, and there is **no key-value reuse across domains**, because two towers that never shared a basis produce keys and values that are not commensurable: a cached tensor is the output of a learned projection and means nothing to a projection trained separately. Domain towers are peers.

On the **size axis** the arrangement is a distillation ladder, one parent per domain with smaller tiers distilled from it. Lineage is shared by construction, and **key-value reuse works along it**. The enabling fact is that the grouped-query attention pattern published families use holds `head_dim` and `num_kv_heads` fixed across sizes while varying only the query-head count and the layer count, so the per-layer key-value tensors of a small tier and a large tier are already the same shape. At the common configuration of 8 key-value heads and a head dimension of 128, one layer’s key-value state for one token is $2 \times 8 \times 128 = 2048$ elements at every tier in the family. What actually blocks reuse is therefore not shape. It is layer-count mismatch, which is a correspondence problem, and the fact that the two tiers’ projections were learned separately even from a common parent, which is a training problem. Both are attackable by training the ladder to agree; neither is a structural impossibility the way the domain-axis case is. The specific family configuration is stated from memory of published model cards and we have not re-read them here; it belongs in `logos/PRIOR_ART_v03.md` for verification, and the argument degrades to “shapes may need aligning too” if it turns out to be false.

The size axis is close to free on the two quantities Proposition 1 says are binding. A 7B and a 70B tier alongside a 2.8T parent add 7.7×10^{10} parameters, which is **2.75 percent** of the tower, or 40.9 GB at MXFP4’s true 4.25 bits, well under one accelerator. Training them by distillation at $D = 20N$ costs $6 \times 7 \times 10^{10} \times 1.4 \times 10^{12} = 5.88 \times 10^{23}$ FLOPs for the 70B tier and 5.88×10^{21} for the 7B, a total of 5.9×10^{23} , which against the tower’s own 2.2 to 3.1×10^{25} is 1.9 to 2.7 percent. Both costs land near the same small number, which is a coincidence of $D = 20N$ and worth no more weight than that.

The “four towers, not five” verdict survives this unchanged. The criterion above is stated entirely on the domain axis: corpus disjointness, objective conflict and update cadence are all properties of what a tower is trained on and for, not of how large it is. Mathematics and Logic still fail all three and still merge. What the size axis changes is the number of served artifacts, not the number of domains, and §11.2 notes the consequence for evaluation: a two-dimensional partition multiplies the arm count, and any metric that takes a maximum over arms degrades as it grows. The full specification of the ladder, including the distillation schedule and the reuse protocol, is in `logos/LADDER_ARCHITECTURE.md`; we do not import it here, because the paper’s claim is only that the second axis exists, is cheap, and does not disturb the first.

4 The router

Every query passes through one router, and a router that dispatches between frontier-grade specialists is where the risk of §3 concentrates. This section gives it a design.

4.1 Where the training signal comes from

A router needs examples of the form (query, tower, outcome). Collecting them by serving every query on every tower costs k times inference, which is affordable on an evaluation set and not in production. The standard treatment is bandit feedback: the router picks one tower, observes only that tower’s reward, and updates only the arm it pulled [7]. Contextual-bandit and cost-aware formulations are established [7, 57], and we adopt them rather than inventing anything.

The design that follows has three parts.

- (i) **A capability prior.** Initialise from per-tower evaluation vectors, so that a fresh router starts at the quality of static domain-label routing rather than at random.
- (ii) **Exploration on a sampled slice.** Route a small fraction ε of traffic off-policy and record the outcome. This is the only source of counterfactual signal, and its cost is ε times the quality gap, which is budgetable.
- (iii) **Outcome signals, where they are free.** In verifiable domains the outcome costs nothing: tests pass or fail, a proof checks or does not, a program runs or crashes.

Part (iii) has a consequence worth stating plainly, because it is a structural property of the architecture and not a tuning detail. **The router will be well trained exactly where verification is cheap, and poorly trained everywhere else.** Code and mathematics generate free labels. Life sciences, policy, and open-ended reasoning do not. So routing quality will be uneven across towers in a way that correlates with domain verifiability, and the towers most in need of good routing are the ones that will get the worst of it. Any deployment should measure routing quality per domain rather than in aggregate, or this asymmetry will hide inside a good average.

4.2 What happens when a tower is unreachable

A distributed five-tower system will lose towers. The contract we adopt is a three-step ladder, and the important property is that every step is visible to the caller:

- (1) Serve from the preferred tower.
- (2) If unavailable, serve from the next-best tower and **attach a degradation flag to the response.**
- (3) If no capable tower is available, hand to the Administration tower, which declines and says why.

Step 2’s flag is the whole point. The Fable-to-Opus demotion path shows what silent degradation costs: a caller who cannot tell that a weaker model answered cannot compensate, retry, or escalate. A degradation flag turns an invisible quality loss into a handled error.

4.3 Router drift when a tower is swapped

Branch-Adapt-Route’s economic case is that towers can be replaced independently. But every replacement invalidates whatever the router learned about the old tower, and if re-learning is expensive then update cost is not linear after all and the architecture’s main advantage evaporates.

The fix is to make the router a function of *tower descriptors* rather than tower indices. Represent each tower by its evaluation vector, train the router over descriptors, and a swapped tower enters as a new descriptor with a warm start from its own evaluation rather than from scratch. Whether the warm start actually recovers most of the routing quality is measurable at the one-billion-parameter scale, and it is the single cheapest test of Branch-Adapt-Route’s economic claim that exists. It appears in §16.

4.4 Users steering the router

If a caller can influence which tower serves them, they can steer around the Administration tower, which would make the compliance mechanism of §13 decorative.

The defence is placement. The safety mask of Eq. (21) is applied *after* routing, as a mask over the permitted expert set, not as a preference the router weighs. A router preference can be outbid by prompt content; a post-router mask cannot, because by the time it applies, the routing decision is already made and the mask simply removes options. This costs nothing and it closes the obvious attack.

One interaction remains open and we flag it as a likely bug rather than a solved problem. Masking experts to $-\infty$ changes the routing distribution, and the quantile balancer of §5.3 balances exactly that distribution across a batch. Masking for one user therefore shifts quantile thresholds for every other token in the batch. A compliance control that quietly degrades unrelated users’ routing is a governance failure and not merely a performance one. Nobody appears to have looked at this. It is falsifier F6.

4.5 This section specifies a known component, and its related work is incomplete

We flag a gap here rather than assert a result, because the check that would close it has not been run. The master router as specified above is a routing layer over heterogeneous models, and that is deployed prior art rather than a new object. OpenRouter operates one commercially: a single endpoint over models from many providers, with an automatic router that classifies each prompt into one of roughly thirty task types and selects a model, reporting back which one served the request [53]. There is also an academic literature on the same problem, of which we verified three entries against their primary sources. RouteLLM trains routers on preference data to dispatch between a stronger and a weaker model at a chosen cost-quality point [50]. FrugalGPT gives the cascade formulation, calling models in sequence until a scored answer is good enough [13]. Hybrid LLM routes by predicted query difficulty against a quality target [18]. These sit alongside the bandit-feedback and survey work this section already builds on [7, 57].

The gap is one of process and we name it as such. The independent round-2 audit of this paper examined the observation-bound literature of §12 for missing related work and found some, and it never examined the routing literature at all. So the four entries above are what we confirmed, not the result of a sweep, and the nearest prior art to the specific design given above may well be closer than anything cited here.

The consequence is a framing correction that applies to the whole paper. §4 *specifies* a known component, with the verifiability asymmetry, the descriptor warm start and the post-routing mask as engineering choices within it, and it should not be read as a contribution. The paper’s contributions sit elsewhere: the four defects diagnosed in §11.2 in a routing metric of our own, whose replacements are adopted from published work and claimed as no contribution at all, the

observation bound of §12 together with the conjecture §12 states against it, and the regulatory-arbitrage argument of Conjecture 3. Nothing about the router is claimed as novel anywhere in this paper, and if any sentence elsewhere reads that way, this one governs.

5 Micro-sparsity inside a tower

5.1 Heterogeneous experts, and why routers abuse them

Tokens differ enormously in the computation they deserve. Punctuation needs almost none; multi-step deduction needs a great deal. Experts of uniform size therefore misallocate by construction. Heterogeneous mixture-of-experts assigns experts of differing sizes, pairing large reasoning experts with light syntactic sub-networks, and adds a training objective that pushes activation toward the smaller ones [74].

That objective is necessary because of a specific pathology. A softmax router minimising cross-entropy learns early that larger experts cut loss faster, and concentrates on them. Smaller experts then get too little gradient to become useful, and the large experts become the throughput bottleneck. We call this the over-activation trap, and heterogeneous mixture-of-experts counters it with a penalty that prices expert capacity into the routing decision [74].

Mixture of Heterogeneous Grouped Experts [36] handles the systems consequence, which is uneven accelerator use and wasted parameters. It groups experts so that size is uniform within a group and varies across groups, adds two-level routing for resource-aware combinations, and uses a group-wise auxiliary loss to steer tokens toward the cheapest group that can handle them.

Remark 1 (A term we could not find). *Informal accounts of this architecture refer to an “All-size Group-decoupling Allocation” mechanism. It does not appear in [74], in [36], or anywhere else we could find. We drop it rather than cite it to a source that does not contain it.*

5.2 Routing in a compressed subspace

Dispatching high-dimensional token vectors across accelerators creates an all-to-all communication bottleneck, and in latency-sensitive serving the dominant cost is moving expert weights out of high-bandwidth memory. That regime is memory-bound, so arithmetic intensity does not matter and bytes moved does.

LatentMoE [45] attacks both costs with one change: a shared down-projection before dispatch and an up-projection after combine, so routing, expert computation, and all inter-device communication happen in a smaller space. A hidden state $\mathbf{x}_t \in \mathbb{R}^d$ becomes $\mathbf{z}_t \in \mathbb{R}^\ell$ with $\ell < d$ under a shared matrix:

$$\mathbf{z}_t = \mathbf{W}_{\text{down}} \mathbf{x}_t, \quad \mathbf{W}_{\text{down}} \in \mathbb{R}^{\ell \times d}. \quad (5)$$

Expert selection is still computed on the full dimension d , which preserves gating expressiveness, while expert computation runs in the compressed space:

$$\mathbf{y}_t = \sum_{j \in \text{top-}k} g_j(\mathbf{x}_t) E_j(\mathbf{z}_t), \quad (6)$$

with the result projected back through $\mathbf{W}_{\text{up}}$. Communication volume and weight-load bandwidth both fall by the compression ratio d/ℓ ; a representative setting compresses 4096 to 1024 [45]. Kimi K3 uses a stabilised variant as its core sparsity framework [40].

Remark 2 (Attributing the speedup correctly). *The reported 3.5-fold inference speedup at matched accuracy is an end-to-end system figure for a deployed model, not an isolated measurement of the projection, and the accompanying parameter-savings framing is a vendor comparison against a chosen baseline [45]. We treat the compression ratio d/ℓ as the architectural quantity and the speedup as Tier B. The projections cost real arithmetic, $4d\ell$ per token per layer, $2d\ell$ for each of the down- and up-projections, which has to be charged against the bandwidth saved. Both must be charged. At $d = 4096$ and $\ell = 1024$ that is 16.8 MFLOP per token per layer, about one percent of a layer’s compute at the active-parameter counts of §2.2 rather than half a percent. The trade is favourable because the regime is memory-bound, and it reverses in a compute-bound one. Finding the crossover needs a measured roofline on real hardware, not an analytic sketch.*

5.3 Load balancing without an auxiliary loss

Conventional mixture-of-experts training prevents routing collapse with an auxiliary load-balancing loss, which by construction fights the primary objective. The auxiliary-loss-free line [75] replaces it with a per-expert bias added to routing scores and updated from recent load, which preserves causality and avoids gradient interference.

Kimi K3 reports a further step, **Quantile Balancing**, in which expert allocation comes directly from router-score quantiles: a token activates an expert when its score lands in the top quantile. The vendor describes this as deterministic, free of tuning hyperparameters, and producing even utilisation with no dead experts [40, 29]. The method is not the vendor’s. It was developed by J. Su and described in a personal blog post in February 2026, months before K3’s release [65], and it has been reported independently at 32B total with 5B active, 64 routed experts of which 4 activate, across 32 layers and 326B tokens at about 10^{22} FLOPs, with no auxiliary loss and no loss spikes [51]. Writing $s[i, j]$ for the router logit of token i at expert j , $b[j]$ for the persistent bias, n for expert count and k for the number activated, the alternating quantile solver computes

$$\alpha[i] = \text{quantile}_{1-k/n}(s[i, :] + b[:]) \quad \text{over experts,} \quad (7)$$

$$\beta[j] = \text{quantile}_{1-k/n}(s[:, j] - \alpha[:]) \quad \text{over tokens,} \quad (8)$$

and sets $b[j] \leftarrow -\beta[j]$ under stopped gradients, so balancing never contaminates backpropagation.

Quantile Balancing works on batch statistics, which leaves per-sequence imbalance untouched. That gap has preprint-grade treatment by a different route: Expert Threshold routing maintains per-expert exponential-moving-average thresholds over the token distribution, which balances load without an auxiliary loss and without any dependence on other tokens in the batch, and reports pretraining results to 2.4B parameters on FineWeb-Edu [69]. The mechanism this architecture adopts instead is the causal counterpart, **Causal Dual Bias**, which tracks expert usage through a dual variable β_t maintained by online dual descent within a sequence, and penalises later tokens by

$$\tilde{s}_t = s_t - \beta_t, \quad (9)$$

with β_t proportional to cumulative imbalance. This balances each sequence without breaking autoregressive causality [12].

Remark 3 (Where these two mechanisms actually come from). *Quantile Balancing originates in a personal blog post [65] and is reported in a vendor technical blog [40], a third-party explainer*

[29], and an independent at-scale training report [51]. Causal Dual Bias is documented in a personal blog post [12]. Neither mechanism has a peer-reviewed or preprint source at the time of writing, though the causal per-sequence problem Causal Dual Bias addresses does have one by a different route [69], and the equations above are our reconstruction from those descriptions and the surrounding auxiliary-loss-free literature [75, 1], not transcriptions from a paper. Anyone implementing them should treat them as specifications to validate. That validation is a one-accelerator experiment and it is falsifier F3.

6 Attention and context at a million tokens

6.1 Kimi Delta Attention

Quadratic self-attention fails at million-token context twice over: computation grows as $O(N^2)$, and the key-value cache exhausts accelerator memory. Kimi Delta Attention [70] is a hybrid linear-attention module that extends Gated DeltaNet with a channel-wise diagonal gate in place of scalar decay, so each feature dimension keeps its own forgetting rate. Unlike uniform-decay linear models it can hold a reasoning state in some channels while clearing others quickly. Its chunkwise algorithm uses a specialised variant of Diagonal-Plus-Low-Rank transition matrices, which cuts computation relative to the general form while staying closer to the classical delta rule [70].

The deployed configuration interleaves it with Multi-head Latent Attention in a 3:1 ratio, giving sequential dependency to the linear layers and exact global retrieval to the latent-attention layers. Reported effect: 75% less key-value cache and up to six times higher decoding throughput at one million tokens [70].

6.2 Why recurrence breaks prefix caching

Recurrent layers create a systems problem that quadratic attention does not have. Paged attention engines cache key-value vectors in predictable physical blocks, so any prefix boundary is a valid resume point. A recurrent layer instead carries a dense running state that updates at every token. Storing it at every possible match boundary is too expensive, and storing it only at coarse boundaries means two requests sharing nearly an entire prompt can miss the reusable prefix [73].

Sparse prefix caching [60] formalises the trade: keep exact recurrent states at a sparse set of checkpoints and, on a hit, resume from the deepest stored checkpoint and recompute the rest of the suffix exactly. Checkpoint placement under a distribution over overlap depths has an exact $O(NM)$ dynamic program, outputs are bit-exact, and no new recurrent kernels are needed.

The production integration separates three quantities that used to be one [73]: the **physical block size** at which state is allocated on device, the **scheduler alignment** where execution must stop so cache groups stay consistent, and the **prefix-match unit** at which a shared prefix is hashed. Separating them allows fine-grained matching inside a larger physical state block. The scheduler stops at the right block and hash boundaries so the recurrent state genuinely corresponds to the advertised prefix, and on a match the cached state is copied into a private destination before the request extends it. Full-attention and recurrent caches keep a synchronised computed-token count despite different physical block sizes.

We spend space on this because it is the clearest case in the architecture of a modelling decision creating a scheduler problem. Choosing channel-wise linear attention for its memory

behaviour produces a cache-coherence problem with no analogue in what it replaced, and that cost is rarely priced when the mechanism is selected.

6.3 Delta Attention Residuals

In very deep networks, PreNorm residual streams dilute early layers: hidden states accumulate, and the running sum swamps individual sublayer outputs. Attention Residuals [31] replace additive residuals with learned softmax attention over previous layer outputs, giving selective cross-layer routing. But attending over *cumulative* hidden states, which are highly similar between adjacent layers, causes routing collapse deeper in the network: attention weights flatten toward uniform, with a maximum around 0.2 [35].

Delta Attention Residuals [35] attend over the change each sublayer contributes,

$$\mathbf{v}_i = \mathbf{h}_{i+1} - \mathbf{h}_i, \quad (10)$$

combined additively so the baseline residual stream survives intact:

$$\hat{\mathbf{h}}_l = \tilde{\mathbf{h}}_l + \sum_i \alpha_{i \rightarrow l} \mathbf{v}_i. \quad (11)$$

Because these increments are structurally diverse and sit in different abstraction subspaces, the attention distribution sharpens, with maximum weight around 0.6. Across scales from 220M to 7.6B, the mechanism beats both standard residuals and cumulative Attention Residuals, with validation-perplexity gains of 1.7 to 8.2 percent [35]. The additive form matters for deployment: at initialisation a uniform softmax output is a bounded perturbation on a preserved residual stream, so existing checkpoints can be converted by finetuning rather than retrained.

The claim worth testing independently is the conversion, not the routing contrast. Measuring that a delta-attention implementation produces delta-attention statistics is circular; measuring whether a real checkpoint converts without destabilising is not. The published gains stop at 7.6B and this architecture applies the mechanism at 2.8T. That is falsifier F4.

7 Quantization

7.1 The memory arithmetic

A 14T ensemble at FP16 needs $2 \times 1.4 \times 10^{13} = 28$ TB of high-bandwidth memory. Four-bit microscaling is not four bits per parameter: MXFP4 stores E2M1 elements in blocks of 32 under a shared E8M0 scale (§7.2), and that scale is worth $8/32 = 0.25$ additional bits per parameter, so true density is 4.25 bits. The footprint falls to $1.4 \times 10^{13} \times 4.25/8 = 7.4$ TB, which a multi-node cluster can hold. NVFP4’s finer blocks of 16 cost 4.5 bits and 7.9 TB. The field, this paper’s earlier version included, routinely quotes the flat four-bit figures, which understate by 6.25 percent for MXFP4 and 12.5 percent for NVFP4; we use the true densities and note that even they are floors, because embeddings and output layers are normally kept at higher precision. On the same basis Kimi K3’s MXFP4 weights occupy about 1.49 TB rather than the widely quoted 1.4 TB. Proposition 1 makes this the constraint sparsity does not touch, because every parameter must be resident whether or not it activates.

7.2 Closing the MXFP4 gap in software

The Open Compute Project’s MXFP4 microscaling format uses blocks of 32 with a shared power-of-two scale, which is cheap in hardware. NVIDIA’s NVFP4 uses blocks of 16 with finer scales and is more accurate, with a reported end-to-end accuracy gap of about ten percent [15].

Two software techniques close it [15]. **Overflow-Aware Scaling** increases the effective dynamic range under the same power-of-two constraint by mapping the block maximum into an overflow-aware range, which cuts error on the low-magnitude bulk of the distribution. **Macro Block Scaling** allocates a higher-precision scale at coarser granularity, an outer multiplier that shifts the input distribution into MXFP4’s representable range before block quantization. Quantization fidelity is destroyed disproportionately by well under one percent of elements, which is why an outlier-targeted second scale buys so much. Together the two reduce the gap from about ten percent to under one percent on average, at a mean GEMM overhead of 6.2 percent.

Remark 4 (Two numbers we corrected and one we wrongly discarded). *Earlier statements of this architecture attributed a nine percent compute overhead to this mechanism. The published figure is 6.2 percent [15], and the nine percent appears to be a conflation with the projection overhead of §5.2, a different mechanism in a different part of the network. Separately, a twelve percent tensor-core die-area saving is sometimes quoted for MXFP4. It is the closing clause of the abstract of [15], the paper cited twice in this remark, which reports “12% relative area savings in tensor cores” for MX against NVFP4. No argument here depends on it.*

7.3 Quantization-aware distillation

Post-training quantization fails on large mixture-of-experts models because calibration cannot resolve routing artifacts. Quantization-aware *training* is worse for a model that has been through multi-stage fine-tuning, reinforcement learning, and merging: applying next-token cross-entropy to the low-precision student forces it to relearn base objectives, overwriting alignment installed later.

Quantization-aware distillation [46] aligns the quantized student to the full-precision teacher’s output distribution under KL divergence,

$$\mathcal{L}_{\text{QAD}} = D_{\text{KL}}(p_{\text{teacher}} \parallel p_{\text{student}}) = \sum_y p_{\text{teacher}}(y \mid x) \log \frac{p_{\text{teacher}}(y \mid x)}{p_{\text{student}}(y \mid x)}, \quad (12)$$

and is reported as stable for exactly the multi-stage post-trained models where quantization-aware training is not, and robust to incomplete data coverage, so recovery does not need the original training data [46]. For a tower ensemble that last property decides the matter, since tower post-training corpora may belong to different parties and distillation needs only teacher logits.

8 Multi-tenancy

8.1 Adapters over a frozen backbone

No operator can host a dedicated 14T cluster per client. The pattern is to freeze the backbone and map client data in through low-rank adaptation [28], with thousands of adapters batched

over frozen experts by segmented gather kernels [14, 59]. Mixture-of-LoRA work [78] treats the adapters themselves as a routed expert set.

8.2 Dimensional collapse under a shared codebook

Discretising continuous states against a shared residual-quantized codebook [32, 55] creates gradient interference under multi-tenant updates: one client optimising the codebook for financial telemetry and another for genomic structure pull the same hierarchical indices in opposite directions.

The resulting failure is a documented mathematical pathology, not a tuning problem. Spectral analysis of vector-quantized autoencoders shows that although the embedding space may be defined at 128 or 256 dimensions, quantization suppresses low-variance components and the model compresses into an effective subspace of roughly **4 to 10 dimensions** [44]. Because this is a structural bias of the commitment objective, rank regularisation forces artificial variance at a cost in loss, and continuous-first warm-up strategies [6] delay rather than prevent it. Performance improves and then degrades as effective dimension rises, with the optimum low [44].

Divide-and-Conquer VQ [44] treats the cause instead of the symptom: split the latent space into several low-dimensional subspaces and quantize each independently. Each subspace then satisfies the model’s preference for low dimensionality, while concatenation expands total expressivity combinatorially without triggering variance suppression. Binding subspaces to tenant-scoped routers also isolates client semantics.

8.3 The weakest link, and where we moved it

Here is an assessment the rest of this section does not support, offered because the reader should have it.

The shared codebook is the least motivated component of the architecture as originally drafted. Residual quantization and semantic identifiers are established for generative retrieval and recommendation [55], where discrete hierarchical identifiers *are* the interface. It does not follow that quantizing continuous hidden states through a shared codebook belongs in a language model’s serving path, and none of the sources cited above put it there. The draft used a real pathology and a real fix to defend a placement neither source endorses.

We therefore restrict the codebook to two positions where discrete identifiers genuinely are the interface: the retrieval and long-term-memory layer, where semantic identifiers index tenant corpora, and the router’s tenant-scoped index. In both, the collapse pathology and its remedy apply as described. In the hidden-state path they answer a question nobody asked. Falsifier F5 tests the mechanism on real hidden states rather than synthetic ones, and §12 produces such states as a by-product, since the observation tokenizer there is a residual-quantized codebook sitting in the position we argue it belongs.

9 Decentralized inference

9.1 The split, and what it does not buy

Running a 14T ensemble in one facility concentrates failure and thermal density. The reference deployment distributes layers across nodes owned by independent parties in the manner of Petals

[9, 43], keeping tokenization, encoding, and primary routing inside the operator’s boundary and dispatching the resulting latent vectors outward for downstream decoding.

Remark 5 (Latents are not confidential). *This split is often justified as protecting proprietary integrity on the grounds that a peer sees only latents. That does not survive contact with the embedding-inversion literature: intermediate representations are routinely invertible to substantial fractions of their input text, and a peer seeing latents across many requests is in a better position than one seeing none. The split does have a real benefit, which is that the router, tokenizer, and codebook are what an adversary would need to replicate the system, and they stay inside. It has essentially no confidentiality benefit for user content. We adopt the split for the first reason and withdraw the second. Any deployment handling regulated data must treat the peer network as an untrusted processor with full visibility, which is a data-protection posture rather than a cryptographic one.*

9.2 Dispatch granularity, and why the Petals topology does not transfer

Two distribution topologies get conflated whenever “Petals-style” is used as shorthand, and they are not interchangeable. What separates them is dispatch granularity, and on a wide-area network it is worth about three orders of magnitude.

A **monolithic sparse mixture-of-experts** routes each token to 16 of 896 experts *per layer*. DeepSeek-V3’s published configuration puts 58 mixture-of-experts layers in a 61-layer stack, and taking a count of that order for a 2.8T model gives roughly $16 \times 60 = 960$ expert dispatches per token. That layer count is an assumption rather than a figure, because Moonshot publishes none for K3, but the conclusion is insensitive to it within a factor of two. Which experts fire is data-dependent, so peer load cannot be predicted from the request, and the resulting traffic pattern is all-to-all rather than pipeline-parallel.

Two distinct quantities follow and they must not be conflated. The 960 figure is a **bandwidth and load** count: that many expert invocations are placed per token. **Latency** is smaller, because dispatch inside a layer is parallel while layers are serial, so a peer round trip is paid once per layer and the serial depth is 60 round trips per token, not 960. At 10 ms per round trip that is 0.60 s per token, or 1.67 tokens per second; at 30 ms it is 1.80 s per token, or 0.56 tokens per second. Neither is an interactive serving rate.

The insensitivity above is worth stating with the sensitivity that sits next to it, because only one of the two is disclosed anywhere else and it is the less consequential one. The conclusion is highly sensitive to the round-trip time. Taking ten tokens per second as an interactive floor, sixty serial round trips per token clears it only when the round trip exceeds $100/60 \approx 1.7$ ms, and measured intra-datacentre round trips for exactly this workload are two orders below that: NVSHMEM reports 24.3 microseconds for mixture-of-experts all-to-all with IBGDA, 18.0 with a CPU proxy and 16.7 device-initiated [47], at which sixty serial round trips is about 1.46 ms per token, or roughly 690 tokens per second. **So the argument below is an argument about wide-area networks and not about sparse dispatch as such**, and a reader is entitled to see how narrow that scope is rather than infer it.

The like-for-like comparison is per response and in serial round trips on both sides. A tower ensemble routes a whole query to one tower: one dispatch, one round trip, then local decode, independent of response length. A peer-sharded monolith pays its 60 round trips for *every* token, so a 500-token response costs $60 \times 500 = 30,000$ serial round trips against one. Reading “960 against one” as a latency ratio conflates invocation count with serial depth; separating the two makes the comparison more conservative per token and considerably wider per response.

A **Mixture-of-Towers** routes a whole query to one tower. That is *one* dispatch per query. The round trip is then paid once at admission and decoding runs local to the tower’s own pool *provided that pool is one low-latency locality*, and that proviso is a deployment assumption rather than a consequence of the architecture. It has to be stated as a premise, because §3 makes each tower internally a heterogeneous mixture-of-experts and §9 makes the reference deployment a peer network: if a tower’s own experts are dispatched to across the same wide-area pool, the tower pays the monolith’s sixty serial round trips too and the gap is zero. The comparison below therefore holds when each tower’s serving pool is a single locality and the monolith’s expert set is not. §11’s case 3 is where the premise is least comfortable, because it puts 60 percent of traffic on a 192-accelerator pool that is itself 60 percent of the monolith’s footprint.

Sixty serial round trips per token against one per query is the difference that decides serviceability, and nine hundred and sixty expert invocations per token against one is the difference that decides bandwidth.

Petals [9, 43] distributes *dense* transformer layers pipeline-parallel, which sits at the coarse end of this axis and is why its reported behaviour over the open Internet is encouraging. Adopting that topology for sparse expert dispatch without re-deriving its cost model is unjustified, and that is the standing architecture-review finding F-13 in `logos/ARCHITECTURE_REVIEW.md`, which is a review finding and not the falsifier F13 of §16. The arithmetic above narrows it and does not close it. Nobody has measured sparse-dispatch cost on real hardware at this scale, and an arithmetic bound is not a measurement.

The components such a deployment would be assembled from are named here as *specified*, and nothing in this paper has been run on any of them. For the pooling role, a worker pool that scales to zero across hardware nobody standardised, there is a published open-source runtime: Beta9 [8], AGPL-3.0, from Beam Cloud, whose distributed configuration provides scale-to-zero serverless workloads, agent-backed externally attached worker pools, and a Tailscale mesh for reaching them. It fills that role alongside FaaS_{MoE} [22], which is the mechanism from the literature rather than a runtime and is evaluated at 2.7B. For the sharding role there is Petals. One gap belongs in the same breath as the components: a Petals-style deployment of a 2.8T mixture-of-experts additionally needs the model files converted into the shard format such a network serves, and nobody has performed that conversion for a model of this class. The sharding half of the stack is therefore not merely unmeasured, it is unbuilt.

9.3 Pinning the cache

Moving key-value state dominates distributed autoregressive inference. Every dialogue turn extends a state that must persist, and moving gigabytes per turn across a peer network is untenable, while expert computation benefits from spreading out. StateFlow [64] separates the two: pin key-value state at a sticky serving site for cross-turn reuse, and optimise expert dispatch and aggregation placement across the network.

Session start anchors the cache at an owner chosen by

$$o_s = \arg \min_{v \in \mathcal{P}(e_s)} \left[\lambda_{\text{lat}} d(e_s, v) + \lambda_{\text{util}} u(v) + \lambda_{\text{mem}} \frac{m_s}{M(v)} \right], \quad (13)$$

minimising over peers $\mathcal{P}(e_s)$ reachable from entry point e_s a weighted sum of network distance, current utilisation, and session memory footprint relative to peer capacity. Intermediate vectors go out to remote experts and results return to o_s for aggregation, so state never migrates between

sites. StateFlow reports more than twice the stable dialogue concurrency of distributed baselines and a 53.0 percent reduction in turn-level p95 latency [64].

Equation (13) is our reconstruction of the selection objective from the reported policy, including the weights. An implementation should prefer the published functional form to ours.

Pinning state has a failure consequence the source does not dwell on and we will: because state deliberately cannot move, losing the sticky owner loses the session. The minimum contract is an explicit session restart with the dialogue history replayed from the client, not a silent hang. §11.3 replaces that minimum with content-addressed snapshots under a named ref, which bounds the loss to a checkpoint interval, and states the constraint that construction has to respect.

9.4 Cache management and elastic experts

At the owner, PiKV [33] treats key-value allocation as an expert-sharded, query-driven process rather than a passive block pool, hashing allocations by time and expert index,

$$s(t, e) = (t \bmod N_{\text{tok}}) \oplus (e \bmod N_{\text{exp}}), \quad (14)$$

and combining compression with activity-based eviction of low-utility tokens. The exclusive-or of two independent moduli collides when N_{tok} and N_{exp} share factors. That is true and misleading in equal measure, because it invites an implementer to reach for coprime moduli. Brute-forcing every pair in $[2, 32]^2$ gives collisions in all 584 coprime pairs and all 377 shared-factor pairs: coprimality is not the fix. The minimal case is $N_{\text{tok}} = 2$, $N_{\text{exp}} = 3$, where two of the four reachable slots each carry two pairs. Nor is the collision forced by slot count, since a disjoint-bitfield construction over the same codomain,

$$s(t, e) = (t \bmod N_{\text{tok}}) + (e \bmod N_{\text{exp}}) \cdot 2^{\lceil \log_2 N_{\text{tok}} \rceil}, \quad (15)$$

is injective. The right statement is that exclusive-or is the wrong mixing operator here regardless of the moduli, and the original instinct, to check the construction against the paper before implementing, stands.

Remote layers run under FaaS_{MoE} [22], which deploys experts as stateless serverless functions with on-demand, scale-to-zero invocation across tenants and configurable expert granularity per function, trading elasticity against invocation overhead. The prototype uses under a third of the resources of a full-model baseline. It is evaluated on a 2.7-billion-parameter model, roughly a thousand times smaller than a tower, and cold start is the obvious failure mode at interactive latency. That is falsifier F7.

9.5 Routing across models, not only inside them

Above the towers sits a Mixture-of-Models problem: how much computation to spend, which models participate, how outputs combine. Static routing on token counts and coarse domain labels ignores within-query difficulty and cross-model complementarity.

A deployed instance is Echo [71], which coordinates open-weight models behind one endpoint and adapts computation per request. The general finding in this line [10, 23] is that complementarity is real: a weaker model can be the right choice on specific sub-problems, so a good router can beat the best single model in its pool. Vendor-reported results place such a system at frontier-comparable quality for about a third of the expected inference cost [71]. That figure is Tier B and self-reported, we found no independent evaluation, and cost-per-quality claims are exactly where the choice of baseline decides the answer.

10 Peer integrity

Distributing activations across independently owned peers admits tampering. Recomputation on trusted hardware defeats the purpose of distributing, and cryptographic commitment over every activation costs too much latency.

The published approach [17] is a threshold-free statistical test using known-answer canaries. The requester picks secret inputs whose correct activations are known in advance and mixes them into ordinary traffic. Peers cannot tell a canary from a live query, so systematic tampering corrupts canaries along with real data.

The difficulty is numerical non-determinism: heterogeneous hardware and non-associative floating-point reduction mean honest peers drift too. The test compares *drift distributions* rather than demanding exact matches. Hardware noise sets a null, and tampering separates from it. Across 408 configurations the detector reaches an area under the ROC curve of 1.0, ranking the malicious shard above every benign shard on every canary [17].

10.1 Why that number will not transfer

Perfect separation across every configuration is a strong result and, read carefully, a narrow one. Three properties of the evaluated setting do not hold here.

- (C1) **The adversary is static.** A shard that tampers on every pass corrupts every canary and is trivially separable. One that tampers with probability $p \ll 1$ corrupts a given canary with probability p , so with m canaries the detection probability is $1 - (1 - p)^m$ and time to detection grows as $1/p$. An adversary who only needs to corrupt a small share of outputs, which is the realistic threat for biasing a code-generation result, can sit below any interactive-latency observation budget.
- (C2) **The numerics are homogeneous.** The benign null is set by hardware drift. A network spanning different accelerators, mixed MXFP4 and NVFP4 paths from §7, and independently compiled kernels has a null that is wider, multi-modal, and node-dependent, which is when a single pooled null swallows the tampering signal. This architecture creates that heterogeneity in §7 and then leans on a detector that assumes it away.
- (C3) **Tampering is assumed to be in value space.** An adversary running a legitimately quantized variant of the assigned expert produces drift that is in-distribution for the benign null while changing the output.

Conjecture 2 (Detection degrades with adversary duty cycle). *Under a fixed canary budget m and an adversary tampering with probability p , detector AUROC falls monotonically in $1/p$, and there exists p^* below which it is statistically indistinguishable from chance at any m compatible with interactive-latency overhead. Widening the benign null to cover heterogeneous numerics raises p^* .*

Testing this needs a benign null *measured* across two accelerators of different models with independently compiled kernels, since a matched pair merely reproduces the published setting. It is falsifier F8, and it is the one we expect to fire against our own architecture.

11 Cost, evaluation, and what breaks

The architecture has argued affordability throughout without pricing anything, and has offered no way to tell a well-routed ensemble from a collection of models behind a lookup table. This section supplies both.

11.1 What it costs to serve

Start from resident memory, since Proposition 1 makes it the constraint sparsity does not move. A 14T ensemble at MXFP4’s true 4.25 bits per parameter is 7.4 TB of weights. On 80 GB accelerators with roughly 70 GB usable for weights, the floor is about **106 accelerators**, before any key-value cache or activation memory, and before the higher-precision embedding and output layers that the floor also ignores.

Real deployments run above that floor, but how far above is not documented and we should not pretend otherwise. A figure of sixty-four or more accelerators for Kimi K3 circulates widely. We could not trace it to a primary document, and the vLLM preview this paper cites for K3 serving does not contain it: it gives a sixteen-accelerator data-parallel and expert-parallel configuration for the MXFP4 mixture-of-experts path and defers hardware requirements to the open-weights release [73]. Taking sixty-four at face value gives 23.2 GB of weights per accelerator, which is 33 percent of a usable 70 GB, 19 percent of a usable 125 GB, and 14 percent of a usable 170 GB. Those figures use K3’s 1.4875 TB at 4.25 bits; dividing a flat-four-bit 1.4 TB instead gives 21.9 GB at 31, 18 and 13 percent, and that substitution is the density error §7 names. The source does not say which part it means, and the reading inverts with the answer: on the largest parts, holding weights to that ratio would mean leaving 86 percent of high-bandwidth memory for cache and activations, which nobody does. Extrapolating the ratio to 14T gives 320 accelerators, or forty eight-accelerator nodes, but that extrapolation assumes the sixty-four is a memory ratio at all, and node granularity is at least as parsimonious an explanation, since $64 = 8 \times 8$ and $320 = 40 \times 8$.

So the statement has to be a bracket: between about 106 accelerators, the memory floor above, and about 320, the ratio-preserving extrapolation. We do not know where in it a real deployment lands, and the gap is key-value cache, activations, expert-parallel communication headroom, and throughput in unknown proportion. Settling it needs a measured serving deployment on named hardware, which is a measurement and not a calculation. The rest of this section uses 320 because it is the conservative end.

Now the comparison the architecture needs and never made. Only one tower activates per query, so a query touches only that tower’s pool. Five towers of 64 accelerators is 320 accelerators. **Five replicas of a single 2.8T generalist of the same per-instance size is also 320 accelerators.** The word is deliberate: a replica fleet built from a domain tower would serve one domain and would not be a competitor to an ensemble at all, so the comparator has to be a model that any query can be admitted to. The hardware bill is the same. The equality is real, and it does not justify the ensemble on traffic entropy alone, because that inference treats the two fleets as interchangeable stock when they differ in the property that decides the answer: **replicas are fungible and towers are not.** Any replica can serve any query, so an idle replica is capacity available to every domain. A Code request cannot consume an idle Life-Sciences accelerator, because the weights resident on it are the wrong weights, and moving 1.49 TB to fix that is not a request-time operation. So the comparison has three cases, not one, and only the third is an argument for the architecture.

Case 1, uniform traffic and uniform sizing. Five domains at 0.20 of traffic each, five pools of 64. Each pool’s share of offered load equals its share of capacity, so both fleets saturate at the same aggregate rate and utilisation is uniform in both. Same cost, same throughput, no serving-cost advantage either way. The ensemble still buys domain coverage, which is a quality argument and is priced nowhere in this section.

Case 2, skewed traffic and uniform sizing. The ensemble is worse. Take shares (0.60, 0.20, 0.10, 0.05, 0.05), a plausible mix for an operator whose users mostly write code. Write C for the aggregate service capacity of the 320 accelerators, so a 64-accelerator pool carries $C/5$. The replica fleet admits any query to any pool and saturates at offered load $L = C$. The ensemble saturates when its busiest pool does, at $p_{\max}L = C/5$, that is

$$L_{\max} = \frac{C}{5 p_{\max}} = \frac{C}{5 \times 0.60} = \frac{C}{3},$$

one third of the replica fleet’s throughput on identical hardware. Mean utilisation at that point is $\sum_i p_i / (5 p_{\max}) = 1/3$, so two thirds of a 320-accelerator fleet is idle while the Code pool queues. This is the loss of statistical multiplexing, it is the standard argument against partitioned capacity, and it applies here in full. **The shares above are stated over the five-way partition of Table 1 and not over the four-way partition §3.5’s criterion returns.** On four towers with Mathematics and Logic merged, shares (0.60, 0.20, 0.15, 0.05), the same skew gives mean utilisation $1/(4 \times 0.60) = 0.417$, that is 58 percent idle and a throughput penalty of 2.4 rather than 3. Both arithmetics are correct for their own partition and neither skew is measured; they must not be interleaved.

Case 3, skewed traffic and asymmetric sizing. The ensemble buys per-domain parameter allocation, and it does not recover case 2’s throughput. A replica fleet is n copies of one model, so every domain gets the same parameter count whether it carries 60 percent of traffic or 5. An ensemble has no such constraint: at a fixed 14T budget, tower sizes can be set proportional to demand. At the same shares that gives towers of 8.4, 2.8, 1.4, 0.7 and 0.7 trillion parameters. At MXFP4’s true 4.25 bits the weights are 4.46, 1.49, 0.74, 0.37 and 0.37 TB, totalling 7.44 TB, which is the same total and therefore the same 106-accelerator memory floor as the uniform arrangement. Allocating the 320 accelerators in the same proportion gives pools of 192, 64, 32, 16, 16, and because both the memory floor and the allocation are proportional to N_i , every pool sits at the same $320/106 = 3.02$ times its own floor. **That equality does not make capacity share equal load share, and it does not recover the multiplexing loss of case 2. The reason is arithmetic rather than judgement.** Memory-floor proportionality is a statement about bytes resident; capacity share is a statement about tokens per second, and the step between them does not hold under this paper’s own conventions. Proposition 1 fixes the sparsity $s = N_{\text{act}}/N_{\text{total}}$ per tower, so tower i activates sN_i parameters, and by Eq. (3) its pool of $A_i = 320 p_i$ accelerators at F FLOP/s serves

$$R_i = \frac{A_i F}{2sN_i} = \frac{320 p_i F}{2s \times 1.4 \times 10^{13} \times p_i} = \frac{320F}{2s \times 1.4 \times 10^{13}},$$

which is *independent of i* . Every pool serves at the same rate whatever its share, the busiest pool still carries 0.60 of the load, and the fleet still saturates at $L_{\max} = C/3$: exactly case 2’s answer, with mean utilisation still $1/3$. The same cancellation survives the memory-bandwidth-bound reading of §5.2, since bytes moved per token in a sparse forward pass is also proportional to activated parameters, and it survives spending the headroom above the floor on replicas of tower i , each serving at a rate proportional to $1/(sN_i)$.

Case 3 would recover the throughput under the opposite convention, activated parameters held *constant* across towers so that the 8.4T tower is three times sparser. That is a legitimate mixture-of-experts design, more experts at unchanged top- k , but it is not the convention Proposition 1 states, it is not what §2.2’s bracket is swept under, and adopting it would change what the case claims: “a tower three times the size” would then buy resident knowledge at unchanged compute per token rather than more compute per query. We keep the fixed- s convention and take the weaker conclusion. **What case 3 does buy is per-domain parameter allocation at throughput parity with case 2:** the 60 percent of traffic that is code is served by a tower three times the size of anything the replica fleet can offer it, and that effect a fixed-size replica fleet structurally cannot reproduce, because to give code queries an 8.4T model it would have to train one and then replicate *it*, at three times the memory bill on every domain including the ones carrying 5 percent. Whether resident parameters per domain is worth a factor of three in throughput is a quality question this section does not price, and the companion logos/LADDER ARCHITECTURE.md §5.3 reaches parity rather than a strict win on its own four-domain instantiation, which is the same answer arrived at independently.

Four caveats bound case 3. The proportionality is exact only because both the memory floor and the allocation scale linearly in N_i ; introduce a per-pool fixed overhead and the small towers stop fitting. The throughput conclusion above is conditional on the fixed- s convention and inverts under constant activated parameters, which is the sensitivity to state. Nothing here says a 0.7T Administration tower is good enough at its job, and §3.5’s criterion does not speak to size. And the traffic mix has to be stable, since re-sizing a tower is a training run and not a scheduling decision, which is the sense in which this ensemble is less elastic than the replica fleet even where it is cheaper.

We spend this much space on the arithmetic because it is checkable and almost nothing else supporting the architecture now is. §2.2 removed the compute motivation, §2.3 removed the token-fragmentation motivation, and §11.4 leaves the data-wall motivation alive only when the residency-bound fraction of each tower’s corpus is below about 0.02, that is only when the shared core is essentially the whole corpus. A cost argument can be checked against an invoice, and being checkable is a property that cuts both ways: the throughput half of case 3 was added and withdrawn inside a single revision of this paper, on arithmetic anyone can redo, which is the behaviour a checkable argument is supposed to have.

One workload is priced here and one is not. §12’s loop is not the serving workload: it shows each observation to $k \geq 2$ towers and admits only the fraction q on which they disagree. All towers are resident in the 320 either way, so the hardware bill does not change and the earlier claim that a two-tower step doubles the footprint would be wrong. What does change is the cost per *admitted* trajectory, which scales as k/q . Neither symbol appears anywhere in this cost model, and §12 notes that q falls by design as the environment is explored, so cost per admitted trajectory rises as the loop succeeds. That is unpriced, and §4’s ε -exploration slice has the same problem: it is charged in quality and never in accelerators.

That yields a decision rule, and it is not the rule this section used to give, nor the rule the previous revision replaced it with. Writing $H(\text{domain})$ for the entropy of the domain distribution over incoming traffic, the oldest rule was that the ensemble’s advantage grows with H and vanishes as $H \rightarrow 0$. That rule is case 1 against case 2 and it holds only under *uniform* tower sizing: an operator whose traffic is eighty percent code and whose towers are all 2.8T should serve code replicas rather than an ensemble. The revision that replaced it said the rule *inverts* under demand-proportional sizing, on the strength of case 3’s now-withdrawn throughput recovery.

What survives the correction is weaker and is stated as such. Demand-proportional sizing does not buy throughput back; it buys *resident parameters where the traffic is*, at the same saturation point case 2 reaches, and a fungible replica fleet has no mechanism for doing so at any H . At maximal H there is nothing to allocate and case 3 collapses into case 1. Reduced to one sentence, *on throughput alone the ensemble never beats a fungible fleet of equal cost, and its serving-side case is that a skewed, stable, pre-known traffic mix lets it spend the same hardware on more parameters per unit of demand*. That is a quality argument wearing a cost argument’s clothes, and this section does not price the quality. It still gives Conjecture 1 something to bite on: the cost of unfiltered frontier deployment is dominated by resident memory per served domain, not by activated compute.

The correction matters beyond this subsection, because the paper’s other motivations have been eroded rather than reinforced. §9.2 supplies a second justification independent of all of them, and it is the sturdier of the two because it does not depend on the traffic mix at all. **A tower ensemble is serviceable over a wide-area heterogeneous pool where a monolithic sparse model of equal total size is not**, because the ensemble pays one peer round trip per query and the monolith pays one per layer. Three things that claim is not. It is not a claim that towers produce better answers, since it is a statement about deployability under a given topology and says nothing about quality. It is not unconditional: it holds under §9.2’s stated premise that each tower’s serving pool is one low-latency locality, which case 3’s 192-accelerator Code pool strains. And it is not measured: §9.2 is arithmetic over an unverified layer count, its crossover sits at a round trip of roughly 1.7 ms so it is an argument about wide-area networks rather than about sparse dispatch as such, and the sparse-dispatch cost model it argues against has never been measured on real hardware either.

For price anchoring, reported blended serving costs per million tokens in mid-2026 are \$2.31 for Kimi K3, \$0.90 for GLM-5.2, and \$0.18 for DeepSeek V4 Pro [38]. All Tier B, and the spread across three trillion-scale sparse models is more than an order of magnitude, which says the dominant term is serving strategy rather than parameter count.

11.2 Measuring whether the router earns its place

The naive metric for an ensemble, the best tower’s score, is exactly what a good router beats and a bad one does not, so it cannot distinguish them. Routing evaluation already has the right primitive: an *oracle ceiling*, the score of a policy that always sends each query to the best model for it, computed from a fully observed reward matrix on an evaluation set [7, 57]. Serving every evaluation query on every tower costs k times inference, which is affordable on an evaluation set and not in production.

With a ceiling and a floor, the tempting quantity is where the router lands between them:

$$\eta = \frac{S_{\text{router}} - S_{\text{best single tower}}}{S_{\text{oracle}} - S_{\text{best single tower}}}. \quad (16)$$

At $\eta = 0$ the router adds nothing over picking one tower and keeping it. At $\eta = 1$ it is said to route perfectly. η is not the primary metric here, and F10’s criterion is not stated in terms of it, for four reasons that are properties of the definition rather than of any deployment.

- (E1) **The denominator can be zero, on a non-negligible event.** If the best single tower already solves every item any tower solves, η is undefined. Three towers at accuracy 0.90, 0.20, 0.20 on a 50-item per-domain slice put that probability at 0.161 over 200,000 Monte

Carlo trials, and the mandate below to report per domain shrinks the slice and raises it further. The two-axis partition of §3.5 shrinks it again.

- (E2) **η is unbounded below.** A router worse than the fixed best tower scores negative with no floor and no stated scale on which to read the magnitude.
- (E3) **The floor is biased high, and the bias grows with tower count.** $S_{\text{best single tower}}$ is a post-hoc maximum over k noisy scores measured on the same evaluation set, so it carries a winner’s curse. With towers of genuinely equal skill at 0.700 on 500 items, the expected maximum computed exactly from the order statistics of the binomial exceeds the truth by 1.16 points at $k = 2$, 2.10 at $k = 4$, 2.90 at $k = 8$ and 3.31 at $k = 12$. That biases η downward for exactly the wide ensembles it exists to judge, and §3.5’s domain-crossed-with-size partition is what pushes k toward the high end of that list.
- (E4) **The ceiling is inflated by chance, and even without chance it is unreachable.** An oracle that counts an item solvable if *any* tower is right counts lucky guesses. On four-choice multiple choice with towers of literally zero skill, $S_{\text{best single tower}}$ is about 0.250 while $S_{\text{oracle}} = 1 - 0.75^k$ reaches 0.763 at $k = 5$, manufacturing a denominator of 0.513 out of independent guessing. The benign case is no better: five genuinely skilled independent towers at $p = 0.70$ give $S_{\text{oracle}} = 1 - 0.3^5 = 0.99757$, so a router at $\eta = 1$ would have to identify, per item, which tower happens to be right. $\eta = 1$ **does not mean “routes perfectly”**. It means “is an oracle”, and a metric whose upper anchor is unattainable by any router at any budget is not a scale anyone can calibrate against.

What to report instead, and it is not ours

The quantity an operator actually chooses between configurations on is a cost-quality curve rather than a normalised gap, and the field already has named, benchmarked formulations of it. **We adopt them. Nothing in this subsection is a contribution, and the correct description of what follows is that the paper stops using a metric it invented and starts using the standard ones.** Proposing a cost-quality metric here would repeat, inside the repair, exactly the missing-related-work defect §12 was corrected for.

The primary metric is **AIQ**, average improvement in quality, from RouterBench [48]:

$$\text{AIQ}(R_\theta) = \frac{1}{c_{\max} - c_{\min}} \int_{c_{\min}}^{c_{\max}} \tilde{R}_\theta \, dc, \quad (17)$$

the area under a router’s cost-quality curve normalised by the span of the cost domain. RouterBench supplies the supporting constructions with it: a non-decreasing convex hull over operating points, which is the Pareto frontier object under its published name, a Zero Router reference built from the individual models’ collective hull, and an Oracle Router [48]. The secondary pair is from RouteLLM [50], which defines the performance gap recovered between a designated weak model M_w and a designated strong model M_s ,

$$\text{PGR}(M_R^\alpha) = \frac{r(M_R^\alpha) - r(M_w)}{r(M_s) - r(M_w)}, \quad \text{APGR} = \int_0^1 \text{PGR}(M_R^\alpha) \, d c(M_R^\alpha), \quad (18)$$

together with $\text{CPT}(x\%)$, the minimum percentage of calls to the strong model needed to reach $\text{PGR} = x\%$. Both directions of the trade were reported before either metric existed, by FrugalGPT, as quality at fixed cost and cost at fixed quality [13]. **We report AIQ as the primary scalar and APGR with $\text{CPT}(50\%)$ as the secondary pair.**

The reason for the swap is specific, and it is the four defects above rather than a general preference.

- (E1) AIQ avoids the zero denominator **outright**: it contains no ratio and no oracle, so a dominating single tower is simply a flat curve with a well-defined area. APGR does *not* avoid it structurally, and we say so rather than claim the pair is clean: $r(M_s) - r(M_w)$ is zero if the designated endpoints score alike. What changes is that the endpoints are designated *a priori* rather than selected post hoc, so the condition is diagnosable before evaluation instead of emerging from per-domain slicing. Weakened, not eliminated.
- (E2) AIQ inherits the quality scale, so with quality normalised to $[0, 1]$ its value lies in $[0, 1]$ and reads as average quality across the cost range. RouterBench does not state that bound as a theorem, and it follows from normalised quality rather than from the paper (logos/PRIOR_ART_v03.md §7, item 11). APGR is stated as bounded in $[0, 1]$, though a single- α PGR can go negative.
- (E3) Neither contains a maximum over k noisy scores. AIQ’s hull is a reference *curve* over the individual models rather than a maximum entering a denominator, so selection bias in the hull moves the comparison line and not the metric value. APGR takes $k = 2$ endpoints fixed in advance.
- (E4) This is the one caveat, and it needs stating precisely. **RouterBench’s Oracle Router is itself a per-item maximum over k and carries the chance inflation in full**, manufacturing $1 - 0.75^k$ of apparent headroom on four-choice items with zero-skill towers, which is the 0.763 at $k = 5$ computed above. The difference that matters is that the oracle is a reported reference *line* and not a term inside Eq. (17), so AIQ’s value is unaffected by it. Report the oracle line only chance-corrected, and say in the caption that it has been corrected. APGR uses no oracle and avoids the defect entirely.

So η is retired rather than kept as a diagnostic. The one thing it was for, the fraction of available headroom a router recovered, is what APGR measures, with endpoints designated in advance instead of maximised over, which is the same idea done correctly by people who published it two years earlier. There is no residual question η answers that APGR does not answer better.

One property matters more than the rest for this paper. AIQ **does not fragment under the two-dimensional partition of §3.5**. Cost and quality are both per-query quantities, so a domain slice or a size slice is read on the same axes, and slices aggregate by summing cost and averaging quality under the traffic weights, which is the operator’s own objective. η is a ratio of differences and does not aggregate at all: a weighted mean of per-domain η is not the η of the pooled set, and with domains crossed by size tiers there is no level at which the ratio is both defined and decision-relevant.

Falsifier F10’s criterion is restated on AIQ, since as written on η it was unmeasurable whenever η was negative or undefined. A warm-started router fails F10 if, on any domain carrying at least five percent of evaluation traffic, its post-swap AIQ falls more than 0.02 below its pre-swap AIQ, or its CPT(50%) rises by more than ten percentage points of calls. Three conditions make that computable. They are not optional: the two metrics are *not* both defined on every evaluation set, the second failing on this subsection’s own showing at (E1).

- (i) **A common cost domain.** Eq. (17) normalises by the span $c_{\max} - c_{\min}$, and a tower swap changes the pool’s price range, so pre- and post-swap AIQ are otherwise averages

over different intervals and a router whose behaviour did not change can move. Both sides are integrated over the *union* of the pre- and post-swap operating ranges, fixed before the swap.

- (ii) **A defined denominator.** CPT(50%) is defined through PGR and PGR is undefined when $r(M_s) = r(M_w)$. The limb is evaluated only on slices where the designated endpoints differ in score, which (E1) already says must be checked in advance; where they do not, F10 is decided on AIQ alone and the exclusion is reported.
- (iii) **A slice size and a multiplicity rule.** 0.02 is below the noise floor of a small slice: at quality near 0.7 on a 50-item slice the standard error of one quality estimate is $\sqrt{0.7 \times 0.3/50} = 0.065$, more than three times the threshold, and with §3.5’s two-dimensional partition there are enough slices that an unchanged router would fail somewhere with high probability. F10 requires a minimum of 5,000 evaluation items per reported slice, a paired test across the swap at a one-sided $\alpha = 0.05$, and a Bonferroni correction across the reported slices; or, alternatively, the 0.02 is replaced by a threshold set from a measured noise floor on the same slices. Without one of those two, F10 is a research task rather than a falsifier.

Per §4, AIQ must be reported per domain. A good pooled AIQ hiding a collapsed one on the unverifiable domains is the expected failure, not a surprise, and it is the failure the verifiability asymmetry of §4 predicts. The design claim this makes falsifiable is the one that matters: if the router’s curve never separates from the best single tower’s point, the ensemble is a collection of models with a lookup table, and the move is to ship the best tower.

11.3 Failure semantics, collected

Three failure paths, one contract each.

Tower unreachable. The three-step ladder of §4, with a degradation flag on the response at step two. No silent substitution.

Sticky owner lost. Session restart with client-side history replay, surfaced as an explicit error. State cannot migrate by design, so pretending otherwise produces a hang. This is the weakest of the three contracts and the next paragraphs replace it with a better one.

Straggler in expert dispatch. All-to-all dispatch sets p99 by the slowest peer. The mitigation is a deadline per dispatch round with a degraded top- k fallback, which lowers quality by a known amount instead of raising latency by an unknown one.

The general principle: at every failure the caller learns that quality dropped. The architecture’s own §1 example of demotion without notification is the pattern to avoid.

Losing the sticky owner should not cost the session

Architecture-review finding F-10 (logos/ARCHITECTURE_REVIEW.md, not the falsifier F10 of §16) states the gap precisely: there is no failure semantics for the sticky key-value owner going away mid-session, because StateFlow pins state so that it never moves, and pinning state means that losing the owner loses the session. Restart-and-replay is a contract for that loss, not an answer to it. Here is an answer.

Separate the state from its name, on the model that content-addressed version control and container registries already use: immutable objects under a content hash, plus small mutable *refs* that name one of them. Git’s object store and its branch refs are the reference implementation; a container tag resolving to an image digest is the same object in a different domain. Applied here, a session’s key-value state is checkpointed as an **immutable, content-addressed snapshot**, and the session is a **named mutable ref** pointing at the most recent one.

Two properties follow, and the second is the one that matters for a peer network. Immutable snapshots replicate freely, because a content hash makes a copy self-verifying and there is no coherence problem between replicas of an object that cannot change. So snapshots can be pushed to several peers and across several failure domains, and the more peers in the pool the more places a given snapshot can live. **Churn resilience therefore improves with pool size rather than degrading with it**, which inverts the usual property of a volunteer network. The ref itself is a pointer of a few tens of bytes and can sit wherever the operator already keeps consistent metadata. Losing the owner then costs the delta accumulated since the last named snapshot, which is bounded by the checkpoint interval, and not the session.

Remark 6 (Key-value state does not merge, and this constraint is load-bearing). *The analogy to version control stops at exactly one place and the design has to stop with it. Text three-way-merges because a document is a sequence of independently meaningful lines with a common ancestor, so two divergent edits can be reconciled by a rule. Attention state is not that. Two consumers who advance the same session’s key-value cache along different continuations hold states that no rule reconciles, and there is no merge function to write. So a session name must advance **fast-forward only, under a single designated writer**. Every other consumer that wants to continue from a snapshot forks: it takes the snapshot, starts its own ref, and diverges by design, which is the intended and cheap operation. Multi-writer access to one name is not a merge conflict, it is a silent-corruption path, because nothing errors. The second writer’s update simply overwrites a pointer, the first writer’s continuation is orphaned, and the surviving state is coherent-looking and wrong. A system that permits it will not produce a diagnosable failure, which is worse than the hang the restart contract was written to avoid.*

One limitation belongs with the answer, because it is what makes the answer specific to this architecture rather than general. Replication rescues the *tower* topology, where a session’s state is an artifact belonging to one tower’s pool and any sufficiently provisioned peer can host it. It does not rescue a single model sharded across peers. There, what would have to survive a peer’s departure is not an artifact but the shard geometry itself, the global assignment of layers or experts to machines, and a global layout is not a thing that can be replicated to a peer the way an object can. This is the same distinction §9.2 draws on latency, arriving from the direction of failure rather than of throughput, and it is one more reason the two topologies should not be discussed as if they were one.

11.4 Overlapping corpora and data residency

Proposition 2 requires tower corpora to overlap, and §13 treats towers as clean residency partitions. Both cannot hold, and the architecture as usually stated contains both without noticing.

The resolution is to split the corpus by residency class rather than by tower. A **shared core** of public, licensed, non-personal data may be consumed by every tower, and Proposition 2’s overlap benefit applies to it in full. **Residency-bound shards** are consumed only by towers certified for that jurisdiction, and contribute to $|\bigcup_i \mathcal{C}_i|$ without being shareable.

The consequence is that the headroom shrinks, and saying only that the amount “depends on how much of the corpus is residency-bound” stops one step short of the number that matters. That is not a quantification, and one line of algebra shows why it matters. Proposition 2’s saving applies only to the shared core, so the effective peak requirement is $\max_i (|\text{core}| + |\text{shard}_i|)$ rather than $\max_i D_{\text{opt}}(N_i)$ over an unconstrained corpus. Write f for the residency-bound fraction of each tower’s corpus. Each tower still needs $|\text{core}| + |\text{shard}_i| \geq D_{\text{opt}}(N_i) = 5.6 \times 10^{13}$, so $|\text{core}| = (1 - f) \times 5.6 \times 10^{13}$ and $|\text{shard}_i| = f \times 5.6 \times 10^{13}$, and the total unique supply the system has to find is

$$(1 - f) D_{\text{opt}}(N_i) + 5f D_{\text{opt}}(N_i) = 5.6 \times 10^{13} (1 + 4f). \quad (19)$$

At $f = 0$ this is Proposition 2’s benefit in full. At $f = 1$ it is 2.8×10^{14} , which is exactly what a 14T monolith needs with no decomposition at all, so the benefit is not merely smaller but zero. The benefit is linear in f , $f = 1$ is reachable rather than pathological for an operator serving only regulated jurisdictions, and nothing in this paper bounds f .

Reading Eq. (19) against the token-supply estimate of §2 is worse than we would like. Against the central unique-stock figure of about 6×10^{13} , the decomposed requirement is already at the estimate at $f = 0$ and exceeds it for f above about 0.02. Against the 90 percent upper bound of 2×10^{14} it survives to $f \approx 0.64$ and fails beyond. So the claim at §2.3 that tower decomposition “converts an impossible unique-data requirement into a possible one” holds only at the optimistic end of a supply estimate and the low end of an unmeasured f . Measuring f is falsifier F12, and it needs no accelerators.

12 The observation bound

Everything so far removes a constraint. Sparsity removes the compute limit. Tower decomposition raises the data ceiling. Four-bit serving handles memory. Adapters handle tenancy. None of it touches the quantity that caps the system once human text runs out, and this section argues that quantity is the real subject of the paper.

12.1 What happens when the tokens are gone

Past exhaustion there are three moves. *Repeat* the corpus, which is bounded: up to about four epochs of repeated data costs almost nothing against unique data, and degrades after [42]. *Synthesise* from existing text, which is derivative, since it resamples a distribution the ensemble already holds. Or *ground*: generate tokens by acting on something and seeing what happens.

The obvious form of the third move is to have the towers talk to each other. That is where the literature has something uncomfortable to say.

12.2 Four results that say the same thing

Four separate lines of work, using different methods on different questions, arrive at one conclusion.

- (1) **Debate is a martingale, and informational diversity does not break it.** Under a Dirichlet-Categorical belief model, standard multi-agent debate induces a martingale over agents’ belief in the correct answer, so expected correctness does not improve over rounds, and the result is proved for homogeneous agents under unweighted belief updates [16]. The

same authors extend it: collective belief remains a martingale with heterogeneous agents. Deliberation reinforces a shared prior rather than correcting it, and holding informatively different views does not by itself change that. Two mechanisms are known to break the martingale, and both are internal to the protocol rather than external to the model. Weighting each agent’s contribution by a calibrated confidence that is positively correlated with correctness turns the belief process into a strict submartingale, so expected correctness strictly increases [80]. Diversity-aware initialisation raises the prior probability that the debate starts from a correct answer, and its own authors are explicit that it does so *without changing the update dynamics*, which remain a martingale [80]. Empirically, multi-agent debate does not reliably beat self-consistency [37].

- (2) **Reinforcement learning with verifiable rewards sharpens rather than expands.** Probing the reasoning boundary with pass@ k at large k , base models catch up with their reinforcement-trained versions across every benchmark and model family tested, and eventually surpass them [76]. The training improves sampling efficiency; it does not enlarge the set of problems that can be solved.
- (3) **Recursive self-training without grounding provably degenerates.** Formalised as a discrete-time dynamical system, self-training exhibits two failure modes: entropy decay, where finite sampling monotonically destroys distributional diversity, and variance amplification, where the absence of persistent grounding produces drift by a random walk. If the fraction of exogenous, externally grounded signal vanishes asymptotically, degeneration follows [77]. Model collapse is the empirical face of this [61], though it is not automatic: collapse is established when synthetic data *replaces* real data, and accumulating both together avoids it [25].
- (4) **Self-play works when something external checks the answer.** The strongest zero-human-data result trains a model to propose and solve its own tasks, and reaches state-of-the-art coding and mathematical reasoning without external data, using *a code executor* to validate proposed tasks and verify answers [81]. The executor is not a detail. It is the entire source of correction.

Silver and Sutton frame the destination [62]: experiential data will come to dwarf human data, with agents learning from continuous streams of action, perception, and feedback. The four results above say something more specific about why, and what the limiting resource is.

12.3 The bound

Putting them together gives a statement we take as the paper’s thesis.

At LOGOS scale, capability is bounded by observation bandwidth, not by parameters or compute. Towers arguing under unweighted belief updates gain nothing, whether or not they hold the same information. Reinforcement learning on a fixed problem set sharpens what is already reachable. Self-training without grounding decays. The only mechanism that reliably adds capability is a channel to something outside the model that can say the model is wrong, and the rate of improvement is the rate at which that channel delivers corrections.

One route around the martingale is documented and it does not escape the bound. Confidence-weighted debate improves correctness only under the assumption that confidence is positively correlated with correctness, and the calibration that buys that correlation is itself purchased with external supervision: reinforcement-learning calibration on a hand-curated harder subset with known answers [80]. The exogenous signal moves from the debate into the calibrator; it is not removed. That is the pattern the bound predicts, and it is the strongest piece of evidence for the bound that the debate literature supplies.

This is a synthesis of established results, not a new claim, and we want to be exact about which parts are ours. The martingale result, the $\text{pass}@k$ boundary result, the degeneration theorem, and the code-executor existence proof all belong to their authors. What we add is three things.

First, the bound is the answer to this paper’s own question. Sections 2 to 11 lift every other constraint on a 10T system, and this one is left standing. An operator who has solved compute, corpus, and memory has a system whose ceiling is set by how fast the world corrects it.

Second, and this is the specific contribution, we state a conjecture that runs against the result we have just cited, and we would rather do that in the open than hide it in a paraphrase. **Towers pretrained on disjoint corpora under different objectives with different alignment histories differ in a way that persona-level heterogeneity does not capture, and that difference is exploitable.** The cited authors say otherwise. Choi et al. extend the martingale explicitly to heterogeneous agents, and Zhu et al.’s diversity result changes only the starting distribution [16, 80]. Both evaluate one base model under different prompts, personas or priors. Our conjecture is that corpus-level difference is a different object from prompt-level difference, because it changes what an agent can condition on rather than how it is asked to respond, and that the martingale extension has not been tested against it. This is Tier C, it contradicts a published theorem’s stated extension, and it is the claim we would most like tested. It is falsifier F13.

The conjecture is stated at $\lambda = 0$ in the sense of §3.3, and the architecture of §3 cannot supply that, because Branch-Adapt-Route needs a common seed. So the object under test here and the object the reference architecture can build are not the same object, and the gap between them is whether a divergent continued-pretraining branch is a large enough perturbation to do what an independent pretraining run would. That is falsifier F14, and until it returns, this subsection should be read as a conjecture about independently pretrained models rather than about the towers of this paper.

We are also careful about what the conjecture would buy if it held. Even a broken martingale is a protocol result about debate. It does not on its own supply the exogenous correction the bound names, which is why the loop below adjudicates against an environment rather than against a majority.

Third, the bound gives an admission rule. If exogenous signal is the scarce resource, trajectories should be kept in proportion to how much of it they carry.

12.4 Where this sits in the literature

Three adjacent positions should be named, because a bound stated as a terminal contribution has to say what it is not. Genewein et al. enumerate four frictions on the path from general to superhuman capability, a data wall, an abstraction barrier, an embodied bottleneck, and multi-agent trust, and explicitly decline to say which of them binds [24]. Ding and Wang add environment determinism as a complementary binding axis, with a bound on chain-task success

and a floor on reward optimisation [19]. Sun et al. state the closest empirical version of our thesis, and they state it nine months earlier: in agentic self-learning, generative-reward-model verification capacity is the main bottleneck, and freezing the verifier induces reward hacking and stalls progress [68].

Our claim is narrower than the first and different from the second. We do not survey the frictions; we name one as the residual after §§2 to 11 lift compute, corpus and memory for one specific 14T architecture, and we instantiate it as a loop that runs on a single accelerator. Against Sun et al. we claim less than the framing above might suggest. Their verification-capacity bottleneck and our observation bound are, as far as we can tell, the same quantity measured on different substrates, and where the two differ is scope rather than content: they measure it inside a search environment with a learned verifier, and we argue it survives when the verifier is replaced by an environment nobody trained. If that difference turns out not to matter, the bound is theirs and the contribution here is the architecture-specific derivation and the loop.

12.5 The loop

Define the yield of a trajectory τ as the surprisal of the observed outcome under the ensemble’s own prediction before it acted:

$$\text{yield}(\tau) = -\log P_M(o_{\text{observed}} \mid \text{context}, a). \quad (20)$$

A trajectory where the towers agreed and were right scores near zero. It is self-confirmation, which is the entropy-decay path of result (3). A trajectory where the towers disagreed and the environment settled it scores high.

The loop, which the companion specification `logos/LOGOS_HARNESS.md` develops as `logos-harness`, runs:

1. Show an observation to at least two towers independently. Each returns a prediction, an action, and its reasoning.
2. Score the divergence between their predictive distributions. Discard candidates where they already agree: a settled disagreement is where the environment’s verdict carries information the ensemble did not already hold, and an agreed prediction that turns out right is self-confirmation.
3. Act on the environment.
4. Let the environment settle the disagreement.
5. Score Eq. (20).
6. Admit the trajectory weighted by yield, and accumulate rather than replace, per result (3).
7. Retrain, and repeat. Disagreement shrinks where the environment has been explored, which pushes generation toward the frontier without being told to.

Steps 2 and 6 do the work. The loop structurally cannot train on its own confident agreement, which is the distribution-narrowing dynamic behind collapse, and step 6’s accumulation is the condition under which collapse is known to be avoidable [25].

Step 2 has a consequence for evaluation that has to be stated, because it bites exactly where this loop meets the forecasting falsifiers of §14. Gating on disagreement selects a subsample on which the ensemble is uncertain, and uncertainty correlates with difficulty, so a Brier or skill score computed on the admitted trajectories is not comparable to a persistence, climatology or market baseline computed over a full question set fixed in advance, and the direction of the bias is not known in advance. Any use of this loop to score forecast skill needs an ungated arm run over the whole roster alongside the gated one. The gate is an admission rule for training data. It is not an evaluation protocol, and using it as one would manufacture a result.

12.6 Two environments, chosen because they fail differently

An emulated game is the cheap test. A Game Boy emulator under a high-throughput reinforcement-learning harness [66] generates trajectories with exact ground truth in memory, and a raw-frames-only benchmark [79] evaluates the trained model in the same terms used for frontier vision-language models. Observations are tokenized by a residual-quantized codebook into a vocabulary shared with text, and the model is a small early-fusion decoder in the Chameleon and Emu3 line [11, 20], with loss masked on observation tokens so it learns to read observations rather than generate them [54]. The decisive experiment holds out a vocabulary: terms scrubbed from the entire text stream, acquirable only through grounded trajectories, probed against a matched control set that does appear in text. Adjudication takes milliseconds, ground truth is exact, and the semantics under test are printed on screen, so the observation channel’s fidelity can be checked pixel by pixel.

Social reality is the test that decides, and the observation channel is the companion psychohistory program’s own pipeline [58]: block states, belief-dispersion observables, community partitions, operator-concentration statistics. Towers propose competing readings of a developing episode with a forecast at a fixed horizon, and reality settles it.

The difference between the two matters more than it looks. An emulator is a deterministic program, and a code executor is too. A sufficiently large model can in principle internalise either, at which point yield legitimately goes to zero and the well runs dry. That is the limit of result (4): the strongest self-play result uses an environment the model could eventually learn to simulate. Social reality cannot be internalised, which is the companion program’s central negative finding about out-of-model agents restated as a resource. Environments that resist simulation are the renewable ones.

Remark 7 (Social reality does not scale as a token source). *A generator producing a few thousand good trajectories per quarter falls further than three orders of magnitude short of relevance to a 2.8T tower; three is the flattering reading. The companion specification recomputes the shortfall from measured quantities of the pipeline named above (logos/LOGOS_HARNESS.md §4.3). The window length and the window count are properties of the analysis rather than estimates; the two remaining inputs, the tokens per trajectory and D_{opt} , are estimates and are named as such below, both running conservative against this paper’s own conclusion. The analysis window is $[\text{onset} - 90 \text{ d}, \text{onset} + 22 \text{ d}]$, so no verdict exists until 22 days of post-onset reality have accrued, and nothing in the design shortens that floor. The verdict is one bit per trajectory, *fires-vs-shuffle*, plus two scalars over the whole window, which is not the dense step-wise signal a training loop consumes. The continuous record the pipeline runs on spans 1764 days, which at that window length is fifteen non-overlapping adjudications on the entire dump, and thirty-nine even at the repository’s looser 45-day separation. And the demonstrated lifetime yield,*

across every substrate and every run the companion programme has performed, is 93 adjudicated windows. A figure of 94 circulates in the companion and does not sum from its own breakdown. Against a 2.8T tower’s $D_{\text{opt}} = 5.6 \times 10^{13}$ tokens, itself the order-of-magnitude placeholder §2.3 declares falsifier F1 against, and pricing a trajectory generously at 10^4 tokens, three thousand trajectories per quarter is 3×10^7 tokens and therefore **6.3 orders of magnitude short**; at the demonstrated lifetime yield it is 7.8 orders. Call it six to eight, not three. The emulator is for volume and reality is for validity; conflating them voids the result. If grounded trajectories help on the emulator but not when reality adjudicates, the mechanism was learning emulator artifacts. Retrospective backtesting against sealed rosters looks like a partial relief at the usual price in hindsight contamination. That price is not payable here: every stand-in tower has read about the episodes on the roster, 82 of 83 roster rows carry onsets before 2025-01-01 so knowledge-cutoff-based model selection leaves one usable row, and entity scrubbing raises the floor without stripping structural recognisability. What it puts in place instead is the block-label shuffle and the matched-calm arm, which generate episodes that never happened and therefore cannot have been memorised.

12.7 What would falsify the bound

The bound predicts an ordering: grounded trajectories beat disagreement-gated self-play, which beats unfiltered self-play, which beats nothing. If unfiltered self-play matches grounded trajectories at matched token counts, the bound is wrong and the observation channel is not the scarce resource. That is falsifier F9, it runs on one consumer accelerator, and it is the cheapest experiment in this paper that can return a decisive negative.

13 Governance

13.1 The threshold

Under the EU AI Act, a general-purpose AI model is presumed to present systemic risk when cumulative training compute exceeds 10^{25} floating-point operations [21]. Meeting the presumption triggers the obligations in the General-Purpose AI Code of Practice [26]: a state-of-the-art safety and security framework, continuous lifecycle risk assessment covering automated cyber-offence, biological and chemical proliferation, and loss of control, a safety and security model report, board-level accountability, and whistleblower protection. Exposure for non-compliance runs to €35 million or seven percent of global annual turnover.

Supervision is national. In the Netherlands the Autoriteit Persoonsgegevens coordinates oversight with attention to transparency and high-risk applications, alongside the Rijksinspectie Digitale Infrastructuur for technical market surveillance [5, 56]. Turning obligations into enforceable mechanism is lifecycle governance work, and the Artificial Intelligence Governance Professional certification [30] is the current credential for it.

13.2 How many models is a Mixture-of-Towers?

The threshold is written against *a model*. A Mixture-of-Towers is five separately pretrained and separately post-trained models joined by a router that is itself trained. Three readings follow, all defensible from the text:

- (R1) **Per-tower.** Each tower is assessed alone. A sparse 2.8T tower at 2.2 to 3.1×10^{25} FLOPs clears 10^{25} by a factor between 2.2 and 3.1. On the bracket of §2.2 that is a comfortable in-scope verdict rather than a marginal one, so the per-tower reading no longer offers an escape and R3 carries correspondingly more of the arbitrage. The margin would still close for a tower an order of magnitude sparser, or trained on a tenth the tokens, so the threshold remains sensitive to design choices a provider controls.
- (R2) **Composed system.** The ensemble is one model whose cumulative training compute is the sum over towers plus the router, about 1.1 to 1.5×10^{26} FLOPs, clearly in scope on any reading.
- (R3) **Router only.** The training the provider performed *as a composition* is the router, whose compute is negligible. On this reading the composed system is an integration of third-party models rather than a model the integrator trained.

R3 is the arbitrage, and it is not obviously wrong. Branch-Adapt-Route’s whole economic proposition is that towers can be sourced, swapped, and upgraded independently [41], and an integrator who obtains five open-weight towers and trains only a router has performed a very small amount of training. Pricing that training is harder than it looks, and this paper does not derive a point figure: §4 specifies the router qualitatively and prices exploration in quality rather than in FLOPs. A bracket is the defensible form. Charging the router’s own signal-gathering at the order of 10^{11} FLOPs per token that a tower forward pass costs, §11.2’s oracle-ceiling evaluation over 10^6 queries of 2000 tokens on five towers is about 10^{21} FLOPs; §4’s exploration slice at $\varepsilon = 0.05$ over 10^{13} production tokens is about 5×10^{22} ; at $\varepsilon = 0.10$ over 10^{14} it is about 10^{24} . The bracket spans three orders and its top end is still ten times under 10^{25} , which makes Conjecture 3 stronger rather than weaker, because the arbitrage survives the most generous accounting we can construct. Whether tower inference spent generating router training signal counts as the integrator’s training compute at all is exactly the accounting question at issue, and whether the presumption attaches to the compute a provider expended or to the compute embodied in what it places on the market is not resolved by the text or the Code of Practice as far as we can determine.

Conjecture 3 (Modular composition is a compute-threshold arbitrage surface). *Any capability threshold defined on cumulative training compute per model is evadable by modular composition, because the compute attributable to the composing step can be made arbitrarily small while the capability of the composed system is bounded below by that of its strongest component. A threshold robust to this must attach to deployed system capability, or to embodied compute including sourced components.*

Naming an arbitrage surface makes it easier to use and easier to close. We judge disclosure correct here because the surface follows immediately from reading the threshold next to any modular-composition paper, and because the parties who need it are regulators rather than victims who can be harmed by it. This is the disclosure calculus the companion psychohistory program applies to its control layer [58], and we keep to it.

13.3 Routing as the compliance mechanism

Residency law forbids sensitive regional data from crossing foreign infrastructure. Modularity is what makes the prohibition enforceable rather than aspirational: the operator restricts domains by region, and when a prompt touching sovereign data would otherwise route through an

uncertified path, a mask forces it elsewhere:

$$G_{\text{masked}}(t)_e = \begin{cases} G(t)_e & \text{if } e \in \mathcal{A}_{\text{safe}}, \\ -\infty & \text{if } e \notin \mathcal{A}_{\text{safe}}. \end{cases} \quad (21)$$

Per §4 the mask applies after routing, which is what makes it non-bypassable. Structurally it is the same object as the classifier demotion of §1: a capability-ordered routing decision, differing only in the objective the ordering encodes.

Supply-chain integrity is enforced before dispatch, with the router hashing the target module against an immutable ledger and aborting on mismatch, so a substituted tower cannot be reached. And per §1, forensic capability must stay local, so that incident analysis cannot be revoked by a third party’s classifier mid-incident.

Remark 8 (What the mask does not guarantee). *Masking the expert set constrains which weights execute. It does not constrain what the permitted weights know, and it does not prevent an unmasked tower from having trained on data a masked tower was forbidden. Residency enforcement at the routing layer controls data in flight, not capability. Treating it as the latter is an error the architecture invites, and §11.4 is where the corpus side of the problem is actually handled.*

14 Relation to the psychohistory program

This paper shares a repository with, and is a deliberate counterpart to, the bounded-psychohistory program [58]. The connection runs in both directions and neither is decorative.

That program’s central negative result is that its forward engine fails against *out-of-model agents*, whose type its hypothesis space never contained and against which more data does not help. It also names *major players*, agents large enough to break mean-field averaging, who must be modelled explicitly. LOGOS-class systems instantiate both. A 10T system executing autonomously over long horizons is a major player by construction, and one that discovers novel exploit chains without source access [52] is out-of-model by demonstration.

The Mixture-of-Towers is also a near-decomposability claim [63] about a machine: dense interaction inside a tower, weak interaction between towers, with an aggregate law over coarse-grained units. It is the same architectural bet that program makes about societies, and it fails the same way, when inter-unit coupling grows and the units stop being independent. Socially that is a critical transition; here it is router-induced correlation, where a router that consistently co-activates towers has rebuilt a monolith with worse ergonomics. Whether the social diagnostic transfers is genuinely open. That program’s measure of independent units is a different estimator on a different substrate, and borrowing the name would not license borrowing the validation, so we record it as a hypothesis rather than a result.

The more useful direction runs the other way. Section 12 needs an observation channel the ensemble cannot fake, and social reality is one, for exactly the reason that program identifies as its own limit. Meanwhile that program’s forward falsifiers are not all blocked on the same thing, and the coupling is narrower than it first appears. Smooth-regime skill is not blocked: it has a one-block pilot over 45 scored steps that beat climatology and tied persistence, which is a measured negative rather than an unstarted test. Of the three that remain unstarted, fixed-point reliability needs a lodged announcement set with outcomes, regime occupancy needs a regime monitor operationalised on a named live series, and only parameter invariance across regime

change needs the coupled engine. All three additionally need a multi-regime social reanalysis corpus that does not exist, which that program’s own ledger names as its structural blocker and as a data-access problem rather than a compute one. The loop of §12 could supply the multi-block forward engine that parameter invariance needs. It does not supply the corpus, the question set, or the monitor. So the dependency is real, narrow, and one-directional, and the earlier framing of “the same missing piece from two sides” overstated it in the direction that flattered both papers.

There is a final convergence, and it is not comfortable for either paper. Psychohistory ends at an observability trilemma: complete predictability would require total observability, which is surveillance and the abolition of independent agents, and that program rejects the limit. This paper ends at an observation bound: capability past the token wall is set by the rate at which the world corrects the model. Both papers therefore terminate on the same quantity, approached from opposite sides. One says you cannot predict people without watching them; the other says you cannot improve past human text without watching something. When the something is people, the two limits are the same limit, and the governance problem of §13 is not a separate section but the place where the argument was always going to end up.

15 What has been measured

Every section before this one is specification. This one is not. It reports the measurements that have been run against a working harness, **logos-harness**, on one 24 GB consumer accelerator, and it exists because a partition criterion that is only argued for is the thing routing was invented to avoid.

The scale caveat comes first and applies to everything below. The models are byte-level decoders of about 1.7×10^7 parameters. Nothing here was measured at 10^{12} parameters, nothing was measured on the five reference domains, and no result below transfers to tower scale by any argument we can make. What these runs can do is kill an operationalisation, which is a different and cheaper service: a criterion that fails on four adjacent programming languages at 17M parameters does not become sound at 2.8T, and the failures reported here are of that kind.

15.1 One criterion, three operationalisations, one survivor

§3.5’s P1 asks whether two candidate towers’ corpora are different enough to justify separate parameters. Three ways of asking it were built and run against the same real corpora — source trees in Python, JavaScript, TypeScript and Markdown, drawn from independent open-source projects. Two of the three are dead.

The null is the whole measurement. Before any of them, one discipline decides whether the numbers mean anything. A separation between two domains is only interpretable against a *within*-domain baseline: how far apart do two samples of the same domain sit? The first baseline built here split each corpus by document, and it was wrong in a way that inflated every result. Documents from one project landed on both sides of the split — shared authorship, house style, licence headers, occasionally copied blocks — so the baseline measured *within-project* redundancy and made every cross-domain gap look larger than it was. Under a project-disjoint split, in which no project appears on both sides, the measured separation of the compression variant fell from +0.0746 to +0.0168. The correction is not a detail; it is the difference between

a positive result and a null one, and it was found only because the baseline was built to be attacked.

P1 as written: token-type Jaccard. Dead. Python against English prose measures $J = 0.3246$, above P1’s cut of 0.30, so the criterion as written merges code and prose into one tower. Moving the cut does not help, because the problem is range rather than placement: four *mutually unrelated* corpora span only 0.26 to 0.32 end to end. The measure has less spread across unrelated domains than the partition needs between related ones, and any threshold placed inside that band is separating noise. The cause is structural — lexical overlap between English-derived technical corpora is dominated by shared identifier and punctuation vocabulary, which is present whatever the corpora are about — so this kills the operationalisation rather than P1.

Normalized compression distance. Dead. Compression distance is model-free, needs no accelerator, and on the first null it appeared to vindicate the criterion, which is why the project-disjoint null was built to attack it. It fails the negative control. Under NCD lower means more similar, and Markdown sits to Python at 0.9661 — closer than JavaScript sits to *itself* across independent projects, at 0.9814. A measure that places two different domains nearer each other than one domain is to a fresh sample of itself cannot support any partition, at any cut point, in either direction. The aggregate separation of +0.0168 is the same fact stated less sharply. Distances are computed with the compressor and dictionary size pinned in the result record, since an unpinned window silently changes what is being measured.

Gradient conflict. Survives. The surviving operationalisation abandons model-free measurement. One model is trained on the mixture of all domains; at checkpoints, the gradient of each domain’s loss is taken with respect to every parameter, and pairs of domains are compared by the cosine between those gradients. The question it asks is the one the architecture actually cares about — *do these domains pull the shared parameters in different directions* — rather than whether their surface statistics differ.

It also removes the failure mode that killed the other two, because **it has no cut point to place**. Within-domain replicates are drawn from *project-disjoint* batches, so the within-domain cosine is the measured spread of “two unrelated codebases in the same language”, and a pair is called separable only when its between-domain cosine falls below the weakest such baseline. The threshold is a property of the data, recomputed at every checkpoint. Nothing about the verdict can be tuned by choosing a number, which is why the two paragraphs above spend their time on range and controls rather than on where 0.30 should have been.

step	within	between	separation	complete
0	0.7963	0.6560	0.1402	no (untrained control)
200	0.3007	0.1308	0.1699	yes
600	0.2548	0.0870	0.1678	yes
1500	0.2105	0.0765	0.1340	yes
3000	0.1595	0.0387	0.1207	yes

Complete separation — every between-domain pair below every within-domain baseline — holds from step 200 through step 3000. The per-pair verdicts are also ones a reader can check

against their own judgement, which neither predecessor managed: prose separates from all three programming languages (Markdown against Python at +0.4808), the programming languages merge with each other, and the pair with the highest agreement is JavaScript against TypeScript at +0.8024, which genuinely should merge.

The untrained control earns its place in the table. At step 0 the measure reports *not* separable, with uniformly high cosines, because early gradients are dominated by generic structure — byte frequency, whitespace, the shape of the loss surface. A single reading at initialisation returns the wrong answer, which is why a trajectory is reported rather than a checkpoint, and it is also the evidence that the measure is not trivially separating whatever it is handed.

What this costs, and what it leaves open. The cost class of F11 has changed and §16 is corrected accordingly. The two operationalisations that needed no accelerator are the two that failed; the survivor requires training a model on the mixture, which puts F11 on the consumer-accelerator list rather than off the hardware ledger entirely. Two things remain open and are flagged rather than extrapolated. Separation narrows monotonically with training, from 0.1699 at step 200 to 0.1207 at step 3000, as both cosines decay; whether it survives a full training run is unmeasured, and a criterion that dissolves at convergence would be no criterion. And the measure has been applied to four programming languages, not to the five reference domains of §3.5. Until it is, the "four towers, not five" verdict remains an argument rather than a measurement, and F11 is not discharged.

15.2 Dispatch traffic, counted rather than derived

§9.2 argues that an ensemble routing a query once is serviceable over a wide-area heterogeneous pool where a monolith paying a peer round trip per layer is not. That argument is arithmetic. It has now also been instrumented: a ledger counts bytes crossing a device boundary from the routing decisions a forward pass actually makes, with home placement by index and no affinity scheduling, which is the pessimistic case for the design being argued for. Against per-layer expert routing, per-segment tower routing moves $3.43\times$ less at 4 routed layers, $10.28\times$ less at 12, and $27.44\times$ less at 32. The ratio is the routed depth, and it improves as towers deepen.

The first version of that measurement was wrong by three orders of magnitude, and the error is worth recording because it is the kind an unchecked derivation invites. A segment dispatch moves the *whole segment*, T token-activations, not a single d_{model} vector; counting it as one vector reported a $1318\times$ advantage where the true figure is the routed-layer count. The corrected ledger is a regression test.

The same instrumentation refuted a claim this paper’s companion specification had been making. It argued that towers buy N -fold parameters at $1\times$ key-value cost and presented that as an advantage over per-layer expert routing. Measured at equal total depth, the two cache *identically*, and the equality is now itself a test. The error was forgetting that expert routing leaves attention dense and shared — only the feed-forward blocks are sparse — so it already buys N -fold feed-forward parameters at $1\times$ cache. What the cache argument does separate towers from is *separately served models*, which share neither a cache nor a trunk. That is a real distinction, and it argues against a router-over-independent-models design rather than against sparse experts. §9.2’s claim is unaffected, because it never rested on cache size.

15.3 Status

Three results, one of them a survivor of two failures, one of them a correction to a claim made elsewhere in this work, and none of them at a scale that settles anything about a ten-trillion-parameter machine. The comparison that would bear directly on §3’s structure — a dense model, per-layer expert routing, and per-segment tower routing trained at matched per-token cost on the same mixture — is running on the same accelerator and is not reported here, because it is not finished. We are not reporting a partial arm ordering: with the dense and expert arms complete and the tower arms still training, any comparison available right now is a comparison of two arms, and the question needs three.

16 Falsifiers

Each entry is a claim, a measurement, and the hardware the measurement needs. One of the fourteen, F11, has now been run in the scoped form of §15 and is partially discharged; the other thirteen have not been run. This table is the complete list of the paper’s *falsifiers*. It is not the complete list of open measurement: the work ledger at `logos/GAPS.md` carries further measurements that are budgeted but carry no falsifier identifier, together with items that are data questions rather than experiments.

An earlier version also stated that two falsifiers, F11 and F12, need no accelerators at all. That is now true of F12 only. Both of F11’s accelerator-free operationalisations were built, run, and failed, and the criterion that survives requires training a model on the domain mixture (§15.1). What changed is the cost class of the measurement, not the wording of the ledger.

Four of these could be discharged by simulation on synthetic inputs, and that route is closed as circular: measuring that an implementation of an equation satisfies that equation is not evidence about anything.

The ledger also has a budget ceiling, and it is worth stating here rather than leaving it to be inferred from the hardware column. F2 is the falsifier that would test the architecture as specified, and its right-hand column says “frontier budget for the real thing”. No such budget is available to this work, and none of the larger configurations the paper reasons about, at 5T, 10T, 100T or 1000T total parameters, is reachable either. What is reachable is the low end of F2’s own ladder: the composition gap $\Delta(N)$ between a modularly composed ensemble and a matched jointly post-trained monolith, measured from about 1B to 70B, which is the programme `logos/LADDER_ARCHITECTURE.md` specifies. That curve is a result whether or not a 14T ensemble is ever built, and §12’s bound is the reason it is not a consolation prize: if observation bandwidth rather than parameter count is what binds at the top of the ladder, then the rungs this work cannot buy are also not the rungs where the interesting quantity lives. The seven falsifiers that fit one consumer accelerator and the two that need none are the rest of the reachable programme, and they are the whole of what can proceed now.

ID	Claim	Falsifying observation	Needs
F1	Sparse towers are not governed by $D_{\text{opt}} = 20N_{\text{total}}$ (§2.3)	A fitted mixture-of-experts scaling law giving compute-optimal D within 20% of $20N_{\text{total}}$ at $N_{\text{total}} \geq 10^{12}$ and 98% sparsity	Scaling sweep

ID	Claim	Falsifying observation	Needs
F2	Branch-Adapt-Route composes at tower scale with router quality high enough that misrouting costs less than monolithic alignment damage	Composed ensemble scoring below a matched jointly post-trained baseline at any scale from 7B upward	8-GPU node (7B to 70B proxy); frontier budget for the real thing
F3	Quantile Balancing gives even utilisation with no dead experts, and Causal Dual Bias removes per-sequence imbalance, without an auxiliary loss (§5.3). The first half is already reported at 10^{22} FLOPs [51], 83 times F3’s budget, so F3 is an independent replication rather than first evidence; a disagreement between the two would locate the effect in the vendor-adjacent configuration rather than in the mechanism	Dead experts under Quantile Balancing, or worse balance and downstream loss than the auxiliary-loss-free bias or than Expert Threshold routing [69], in a real 1B, 64-expert run	One consumer GPU
F4	An existing checkpoint converts to Delta Attention Residuals by finetuning without destabilising	Conversion diverging, or failing to recover baseline perplexity, on a real open checkpoint	One consumer GPU
F5	Shared-codebook quantization collapses to 4 to 10 effective dimensions on real hidden states, and subspace partitioning fixes it	Monolithic quantization retaining near-full effective rank on real states, or partitioning failing to recover it	One consumer GPU (by-product of F9)
F6	The post-routing compliance mask of Eq. (21) does not disturb quantile-balanced routing for other users in the batch	Measurable shift in other tokens’ expert assignment when one request is masked	One 8-GPU node
F7	Scale-to-zero serverless experts are viable at interactive latency for tower-scale experts	Cold-start p95 exceeding the per-token latency budget at tower-scale expert sizes	One 8-GPU node
F8	Canary drift detection holds against a realistic adversary (Conj. 2 says it does not)	AUROC staying above 0.95 as duty cycle $p \rightarrow 0.01$ under a null measured across heterogeneous accelerators	Two different GPUs
F9	Grounded trajectories beat disagreement-gated self-play, which beats unfiltered self-play, at matched token counts, without collapse (§12)	Unfiltered self-play matching grounded trajectories, or the collapse monitor firing on the grounded arm	One consumer GPU

ID	Claim	Falsifying observation	Needs
F10	Router quality survives a tower swap via descriptor warm-start, so update cost stays linear (§4)	On any domain carrying at least 5% of evaluation traffic, the warm-started router’s post-swap AIQ [48] falling more than 0.02 below its pre-swap AIQ, or its CPT(50%) [50] rising by more than ten percentage points of calls (§11.2)	One consumer GPU
F11	The criterion of §3.5 returns <i>four</i> towers on real corpora rather than the five Table 1 prints, because Mathematics and Logic fail all three axes. Asserting the five-way partition here instead, as this row once did, makes its falsifying observation identical to the paper’s own conclusion, so nothing could falsify it. Partially run (§15.1): the lexical operationalisation is dead, its total range across unrelated corpora being narrower than the gap it must resolve, and so is the compression-based replacement, which under a project-disjoint null rated two different domains closer than one domain is to a fresh sample of itself. The criterion that survives is <i>gradient conflict</i> , whose threshold is the measured within-domain spread rather than a number chosen by anyone	Gradient conflict failing to hold complete separation over a full training run, rather than the 3000 steps measured, since separation narrows monotonically over that range ($0.1699 \rightarrow 0.1207$); or, applied to the five reference domains rather than the four programming languages measured so far, the criterion returning five towers after all	One consumer GPU. The “no accelerator” entry this row previously carried is withdrawn: both accelerator-free operationalisations failed, and the survivor requires training a model on the mixture

ID	Claim	Falsifying observation	Needs
F12	Tower decomposition converts an impossible unique-corpus requirement into a possible one under real residency constraints (§11.4)	A measured residency-bound fraction f high enough that Eq. (19)’s $5.6 \times 10^{13}(1 + 4f)$ exceeds the central unique-token supply estimate of §2, 6×10^{13} , that is f above about 0.02. The supply point is named because the verdict is a function of it, and leaving it unstated hides that: at the 90 percent lower bound of 2×10^{13} the condition holds at $f = 0$ and F12 fires unconditionally, and at the upper bound of 2×10^{14} it needs $f > 0.64$. At $f = 1$ the benefit is exactly zero on any supply figure	No accelerator; corpus residency classification against a named jurisdiction set

ID	Claim	Falsifying observation	Needs
F13	Corpus-level difference between towers is exploitable by debate in a way persona-level heterogeneity is not, contradicting the heterogeneous-agent extension of the martingale (§12). The claim has two limbs, and they need different instruments	(a) Debate between models of genuinely different pretraining lineage tracking the martingale as closely as debate between personas of one model. (b) Calibrated-confidence weighting lifting ensemble accuracy without any environment adjudication, <i>once the calibrator’s own exogenous supervision is charged to the same ledger and held identical across arms</i>, which would locate the gain in the protocol rather than in the observation channel. Without that condition the kill condition would be the content of [80]’s Theorem 1, a published result §12 cites and accepts, which would make the limb a restatement rather than a test. A null here is an equivalence claim and is declared only under the two-one-sided-tests procedure, at the margin stated in <code>logos/F9_PREREGISTRATION.md</code> §10	(b) One consumer GPU, an arm of F9, 36.3 GPU-hours at the planning roster (18.1 to 290.1 across candidate proposer rosters). (a) One consumer GPU, but <i>not</i> an arm of F9 and not with 350M stand-ins: existing open-weight models of distinct pretraining lineage, quantized and stepped sequentially, 17.4 GPU-hours at a four-model 7-to-8B-class roster and 2.2 at the 1B-class roster F9’s own generation ledger plans on, so the two figures are not addable without saying which roster each is at. The 17.4 assumes prefix caching across debate rounds; without it the same protocol is 34.2

ID	Claim	Falsifying observation	Needs
F14	A common seed with divergent continued pretraining preserves the informative difference F13 needs, so the lineage value Branch-Adapt-Route requires ($\lambda = 1$) and the value F13's conjecture is stated at ($\lambda = 0$) are interchangeable (§3.3)	Debate between two continued-pretraining branches of one base checkpoint tracking the martingale measurably more closely than debate between two independently pre-trained models of matched size and matched benchmark quality, under an identical protocol, where “measurably” means outside the two-one-sided-tests equivalence margin stated in <code>logos/F9_PREREGISTRATION.md</code> §10. If it does, the architecture cannot have both Branch-Adapt-Route and the diversity conjecture, and §12's conjecture is being tested on an object the reference architecture cannot build. F14 is interpretable only conditional on F13 limb (a) first returning a nonzero lineage effect , and the two run in that order: if limb (a) finds no lineage effect at all, both of F14's arms track the martingale equally by construction, the difference is zero, and F14 passes automatically. That pass is not support. It would mean the two lineage settings are interchangeable because neither has any diversity to preserve, and a null F14 under a null F13(a) is reported as uninformative	One consumer GPU, the same instrument as F13 limb (a) and costed with it at 17.4 GPU-hours at the 7-to-8B-class roster. Its binding constraint is not compute but availability: it needs a published continued-pretrain of a base that is also on the proposer roster, and we have not checked that one exists

Seven of the fourteen run on one consumer accelerator: F3, F4, F9 and F10 as independent experiments, with F5 and F13's limb (b) falling out of F9, the latter at 36.3 GPU-hours and

with its own kill condition stated ahead of the run, and F14 sharing an instrument with F13’s limb (a). F13’s limb (a) runs on that accelerator too, but not as an arm of F9 and not on 350M stand-ins. What it needs is models whose pretraining corpora, objectives and alignment histories genuinely differ, and that is a property of how a model was trained rather than of the hardware it runs on. Models with that property already exist, having been pretrained by different organisations on different data under different objectives, and they are arguably a better instrument than five towers from one laboratory, which would share data-collection pipelines and filtering decisions and be less independent than they look. Distinct lineage is the treatment variable: two models from one laboratory, or two finetunes of one base checkpoint, do not count as distinct. So instrumented, limb (a) tests the claim at the level of independently trained open-weight models rather than at tower scale inside one architecture, and the ensemble under test is not a Mixture-of-Towers; that limitation is real and we state it rather than let the limb read as more than it is. Limb (a) is priced at 17.4 GPU-hours, derived from the proposal volume and the roster’s inference cost rather than assumed, at a four-model 7-to-8B-class roster and under an assumption of prefix caching across debate rounds; the same protocol at F9’s own 1B-class generation roster is 2.2 GPU-hours and without prefix caching it is 34.2, so the figure is only meaningful with its roster and its caching assumption attached. F14 runs on the same instrument but is read second, because a null F14 is uninterpretable unless limb (a) has first returned a nonzero lineage effect. The five-tower ensemble itself remains out of reach on any hardware we have, and that is F2, not F13. F12 needs no accelerator; it is a corpus-residency classification. F11 was budgeted the same way and is not: §15.1 reports both of its accelerator-free operationalisations failing, and the criterion that survived them needs a model trained on the domain mixture, which moves F11 onto the consumer-accelerator list and makes eight of the fourteen run there. It remains among the cheapest tests of the paper’s motivation for towers now that §2.2 has removed the compute motivation — the run reported in §15.1 is minutes of one consumer card, not hours — but “cheapest” and “free” are different claims and only the first survives. F9 is the one we would run first, because it tests the constraint that none of the architecture relaxes and it can return a decisive negative on a desktop.

17 Conclusion

A ten-trillion-parameter system is buildable, and not because anyone solved the data wall. The data wall was never the constraint the dense arithmetic claimed. Sparsity separates the compute bill from the parameter count. Modular composition separates the peak unique-corpus requirement from the total parameter budget. Four-bit serving separates the memory footprint from training precision. What is left binding is a shorter and less romantic list: unique high-quality tokens, resident memory, and the quality of a router that must adjudicate between frontier-grade specialists on far less training signal than any of them received.

We have tried to be exact about which joints are load-bearing and untested. The extrapolation of Branch-Adapt-Route from 7B to 2.8T is the central bet and it is unhedged. The shared codebook was in the wrong place and we moved it. The confidentiality story for peer dispatch does not hold and we withdrew it. The peer-integrity result is real, and we predict it will not survive an adversary who is trying. The architecture’s own partition criterion says it wants four towers rather than five, on a domain axis crossed with a size axis that costs 2.75 percent of the parameter budget. Serving five towers costs the same as five replicas of one generalist of the same per-instance size, but the two fleets are not interchangeable, because replicas are fungible

and towers are not: under skewed traffic and uniform sizing the ensemble is strictly worse, losing two thirds of a 320-accelerator fleet to idleness on the five-way partition, or 58 percent of it on the four-way partition §3.5’s criterion actually returns. Sizing towers in proportion to a traffic mix known before they are trained does *not* recover that throughput, though it looks as though it should: under Proposition 1’s fixed sparsity every pool serves at the same rate regardless of its size, so demand-proportional sizing buys resident parameters where the traffic is and leaves the saturation point exactly where the uniform arrangement put it. That is one axis, and §9.2 adds a second that does not depend on the traffic mix: an ensemble routes a query once where a monolith of equal total size pays a peer round trip per layer, so the ensemble stays serviceable over a wide-area heterogeneous pool and the monolith does not, provided each tower’s own pool is one low-latency locality. That axis is arithmetic, unmeasured, conditional on that premise, and scoped to networks whose round trip exceeds about 1.7 ms. The update-economics argument, the alignment-isolation argument and the informative-difference conjecture of §3 are independent of both and are not priced here.

The regulatory finding is the one we would most like to be wrong about. A capability threshold written against *a model* does not survive an architecture whose economic proposition is that models are components. Modular composition makes the compute a provider expends arbitrarily small while the capability it places on the market is bounded below by its strongest sourced part. Closing that gap means attaching the presumption to deployed capability, or to embodied compute, and the sooner it is settled the smaller the incentive to build this architecture specifically to exploit it.

But the finding that changed our own view is the last one. Every mechanism in this paper lifts a constraint, and after all of them are lifted one remains, which none of them touches. Debate among towers is a martingale whether or not they know the same things. Reinforcement learning on a fixed problem set sharpens what was already reachable. Self-training without grounding decays. Self-play works when a code executor checks the answer, and the executor is doing the work. At this scale the limiting resource is not parameters or compute or even data, but observation: the rate at which something outside the model can tell it that it is wrong. Scale buys the capacity to be corrected quickly. It does not buy the corrections.

That is also where the tower architecture stops being an accounting convenience, and where we make our one claim against the literature rather than from it. The debate theorem does not say that improvement needs participants who know different things; it says the opposite, and extends the martingale to heterogeneous agents. We conjecture that towers pretrained on disjoint corpora differ in a way that heterogeneity of prompt or persona does not, and that this difference is exploitable. It is falsifier F13, it is the least defended claim in the paper, and it is the one we would run second. §3.3 makes it worse before it makes it better: Branch-Adapt-Route needs a common seed, so the towers this architecture can actually build are branches of one checkpoint rather than independent models, and whether a branch is a large enough perturbation to carry the conjecture is falsifier F14. Neither of its limbs needs hardware we do not have. The confidence-weighting limb is an arm of F9 at 36.3 GPU-hours, and the diversity limb needs no towers built for the purpose: models of genuinely different pretraining lineage already exist, and the test is to use them, at the price of measuring the claim at the level of independently trained open-weight models rather than at tower scale.

One last thing about status, stated here because the abstract states it and the two should not disagree. The machine this paper specifies is not buildable at any budget behind this work, and the larger configurations it reasons about are further away still. That is not a hedge and it

is not an apology, because it is the same sentence as the paper’s own conclusion: if the binding quantity at the top of the ladder is observation bandwidth rather than parameters, then an unreachable rung is not merely unaffordable, it is the wrong thing to be reaching for. What remains reachable is the part that was going to carry the information anyway. The composition-gap curve $\Delta(N)$ from about 1B to 70B is measurable on hardware that exists here, it is falsifier F2’s own quantity evaluated where it can be evaluated, and it is useful in both directions: it tells an operator with a frontier budget whether modular composition holds at scale, or it tells them not to spend the run. F9 and the other consumer-accelerator falsifiers are reachable. F12 needs no accelerator; F11 was thought not to and does, for the reason §15.1 gives, and it has been run on the hardware that exists here. That is the programme, and none of it waits on the 14T ensemble.

Almost nothing here is a result. It is a specification with its arithmetic checked, its citations audited against primary sources, its weakest components relocated or withdrawn, and fourteen falsifiers stated with the hardware each requires. Eight of them fit on a desktop and one needs no accelerator. The exception is §15, which reports the one falsifier that has been run: two ways of asking F11’s question are dead, a third survives on four programming languages at 1.7×10^7 parameters, and neither the reference domains nor a full training run has been reached. The architecture itself remains untrained and unserved at every scale, and the measurements are instruments for the falsifiers rather than instances of the machine.

AI contribution declaration

This paper was developed in close collaboration with Claude, a large language model from Anthropic. The AI contributed to drafting, to verifying every citation against a primary source, to the arithmetic corrections of §2, and to the synthesis of the four results in §12. The human author, Wingston Sharon, directed the work and contributed its central conjectures: the Mixture-of-Towers composition, the compute-economics reading of restricted deployment (Conjecture 1), the regulatory-arbitrage argument (Conjecture 3), and the proposal that towers can bootstrap past the token wall by generating and adjudicating their own trajectories, which §12 develops. He made all final judgments and takes responsibility for the content, including any errors. An AI system cannot bear accountability for a publication and so cannot be a formal author; this declaration preserves the transparency a joint byline was meant to express.

References

- [1] A theoretical framework for auxiliary-loss-free load balancing of sparse mixture-of-experts in large-scale AI models. *arXiv:2512.03915*, 2025.
- [2] Anthropic. Claude Fable 5 and Claude Mythos 5. <https://www.anthropic.com/news/claude-fable-5-mythos-5>, June 2026.
- [3] Anthropic. Project Glasswing: securing critical software for the AI era. <https://www.anthropic.com/glasswing>, 2026.
- [4] Anthropic. Redeploying Claude Fable 5. <https://www.anthropic.com/news/redeploying-fable-5>, July 2026.
- [5] Autoriteit Persoonsgegevens. Wetgevingstoets Uitvoeringswet AI-verordening. <https://www.autoriteitpersoonsgegevens.nl/>, June 2026.

- [6] Continuous first, discrete later: VQ-VAEs without dimensional collapse. *arXiv:2605.06870*, 2026.
- [7] Learning to route LLMs from bandit feedback. *arXiv:2510.07429*, 2025.
- [8] Beam Cloud. Beta9: ultrafast serverless GPU inference, sandboxes, and background jobs. <https://github.com/beam-cloud/beta9>, AGPL-3.0, 2026. **Named as a specified component. Nothing in this paper has been deployed on it. Scale-to-zero workloads, agent-backed external worker pools and the Tailscale mesh setting are read from the upstream repository’s documentation and default configuration.**
- [9] A. Borzunov et al. Petals: collaborative inference and fine-tuning of large models. *Proceedings of ACL 2023 (System Demonstrations)*, pp. 558–568, 2023.
- [10] Brick: spatial capability routing for the Mixture-of-Models paradigm. *arXiv:2606.13241*, 2026.
- [11] Chameleon Team (Meta FAIR). Chameleon: mixed-modal early-fusion foundation models. *arXiv:2405.09818*, 2024.
- [12] J. Chang. Causal routing bias for aux-loss-free MoE training. <https://jonathanc.net/blog/causal-routing-bias>, 2026. **Blog post; no peer-reviewed source.**
- [13] L. Chen, M. Zaharia, and J. Zou. FrugalGPT: how to use large language models while reducing cost and improving performance. *arXiv:2305.05176*, 2023.
- [14] L. Chen et al. Punica: multi-tenant LoRA serving. *Proceedings of MLSys*, 2024.
- [15] J. Chhugani, G. Jeong, B.-Y. Su, Y. Pan, H. Yang, A. Ankit, J. Yu, S. Deng, Y. Chen, N. Satish, and C. Kim. Unveiling the potential of quantization with MXFP4: strategies for quantization error reduction. *arXiv:2603.08713*, 2026.
- [16] H. K. Choi, X. Zhu, and S. Li. Debate or vote: which yields better decisions in multi-agent large language models? *arXiv:2508.17536*, 2025; *Advances in Neural Information Processing Systems*, 38, 2025 (Spotlight). **Source of the Dirichlet-Categorical martingale result and of its heterogeneous-agent extension.**
- [17] M. Cihangiroglu and A. Nocera. Integrity of peer-to-peer distributed LLM inference under malicious nodes. *arXiv:2607.19490*, 2026.
- [18] D. Ding, A. Mallick, C. Wang, R. Sim, S. Mukherjee, V. Ruhle, L. V. S. Lakshmanan, and A. H. Awadallah. Hybrid LLM: cost-efficient and quality-aware query routing. *International Conference on Learning Representations*, 2024; *arXiv:2404.14618*.
- [19] L. Ding and X. Wang. Grounded scaling: why agentic AI needs deterministic environments. *arXiv:2606.22495*, 2026.
- [20] Emu3 (BAAI). Multimodal learning with next-token prediction for large multimodal models. *Nature*, s41586-025-10041-x, 2025; preprint *arXiv:2409.18869*.
- [21] European Union. Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act), Article 51 and Annex XIII. *Official Journal of the European Union*, 2024.
- [22] FaaSmoE: a serverless framework for multi-tenant mixture-of-experts serving. *arXiv:2604.26881*; *MobiSys Workshops ’26*, 2026.

- [23] Towards generalized routing: model and agent orchestration for adaptive and efficient inference. *arXiv:2509.07571*, 2025.
- [24] T. Genewein, M. Franklin, A. Lerchner, L. Orseau, S. Albanie, A. Bales, C. Wyeth, S. Chan, I. Gabriel, J. Z. Leibo, A. Dafoe, M. Hutter, T. Graepel, and S. Legg. From AGI to ASI. *arXiv:2606.12683*, 2026.
- [25] M. Gerstgrasser et al. Is model collapse inevitable? Breaking the curse of recursion by accumulating real and synthetic data. *arXiv:2404.01413*, 2024.
- [26] European Commission AI Office. General-Purpose AI Code of Practice, final version. <https://code-of-practice.ai/>, 2025.
- [27] J. Hoffmann et al. Training compute-optimal large language models. *arXiv:2203.15556*; *Advances in Neural Information Processing Systems*, 35, 2022.
- [28] E. J. Hu et al. LoRA: low-rank adaptation of large language models. *arXiv:2106.09685*; *ICLR*, 2022.
- [29] K. Huang. Demystifying Kimi K3: the three algorithms behind the top frontend coding model. <https://kenhuangus.substack.com/>, 2026. **Third-party explainer.**
- [30] International Association of Privacy Professionals. Artificial Intelligence Governance Professional (AIGP) certification. <https://iapp.org/certify/aigp>, accessed 2026.
- [31] Kimi Team. Attention Residuals: technical report. *arXiv:2603.15031*, 2026.
- [32] D. Lee et al. Autoregressive image generation using residual quantization. *Proceedings of CVPR*, 2022.
- [33] D. Liu et al. PiKV: KV cache management system for MoE architecture. *arXiv:2508.06526*, 2025.
- [34] J. Ludziewski et al. Joint MoE scaling laws: mixture of experts can be memory efficient. *arXiv:2502.05172*, 2025.
- [35] C. Luo, Z. Cai, and J. Hu. Delta Attention Residuals. *arXiv:2605.18855*, 2026.
- [36] Z. Ma et al. Mixture of Heterogeneous Grouped Experts for language modeling. *arXiv:2604.23108*, 2026.
- [37] A. Smit, P. Duckworth, N. Grinsztajn, T. D. Barrett, and A. Pretorius. Should we be going MAD? A look at multi-agent debate strategies for LLMs. *arXiv:2311.17371*, 2023; *Proceedings of ICML*, 2024.
- [38] Kimi K3 vs DeepSeek V4 Pro vs GLM-5.2: open trillion-scale MoE models compared on benchmarks, license, and serving cost. MarkTechPost, July 2026. **Tier B, secondary.**
- [39] Holistic scaling laws for optimal mixture-of-experts architecture optimization. *arXiv:2603.21862*, 2026.
- [40] Moonshot AI. Kimi K3 technical blog: open frontier intelligence. <https://www.kimi.com/blog/kimi-k3>, 2026. **Vendor deployment report. Not the primary source for Quantile Balancing, which originates with J. Su months earlier; see su2026qb.**
- [41] J. Morrison, S. Adhikesaven, A. Bhagia, M. Zaharia, N. A. Smith, and S. Min. Train separately, merge together: modular post-training with mixture-of-experts. *arXiv:2604.18473*, 2026.

- [42] N. Muennighoff et al. Scaling data-constrained language models. *arXiv:2305.16264*; *Advances in Neural Information Processing Systems*, 36, 2023.
- [43] A. Borzunov et al. Distributed inference and fine-tuning of large language models over the Internet. *Advances in Neural Information Processing Systems*, 36, 2023.
- [44] Dimensional collapse in VQVAEs: evidence and remedies. *Advances in Neural Information Processing Systems*, 38, 2025. Introduces Divide-and-Conquer VQ.
- [45] LatentMoE: toward optimal accuracy per FLOP. *arXiv:2601.18089*, 2026.
- [46] Quantization-aware distillation for NVFP4 inference accuracy recovery. *arXiv:2601.20088*; NVIDIA Nemotron technical report, 2026.
- [47] NVSHMEM device-initiated communication for mixture-of-experts all-to-all dispatch: IBGDA against CPU-proxy round-trip latency. *arXiv:2604.00317*, 2026. **Source of the 16.7 to 24.3 microsecond intra-datacentre round trips quoted in §9.2. Read for the round-trip figures only.**
- [48] Q. J. Hu, J. Bieker, X. Li, N. Jiang, B. Keigwin, G. Ranganath, K. Keutzer, and S. K. Upadhyay. RouterBench: a benchmark for multi-LLM routing systems. *arXiv:2403.12031*, 2024. **Source of AIQ, and of the non-decreasing convex hull used here as the cost-quality frontier. Its Oracle Router is a per-item maximum over k models and carries the chance inflation of §11.2 in full; it is a reported reference line, not a term inside AIQ, so AIQ’s value is unaffected.**
- [49] A. Q. Jiang et al. Mixtral of Experts. *arXiv:2401.04088*, 2024. **Routing analysis (§5) reports no topical specialisation among experts on Pile subsets (ArXiv, PubMed, PhilPapers): assignment tracks syntax, token identity and positional locality, and the distribution over experts is similar across subsets.**
- [50] I. Ong, A. Almahairi, V. Wu, W.-L. Chiang, T. Wu, J. E. Gonzalez, M. W. Kadous, and I. Stoica. RouteLLM: learning to route LLMs with preference data. *arXiv:2406.18665*, 2024.
- [51] L. Dial (Open Athena). Mixture of Experts Quantile Balancing: validated at 32B-A5B (1e22 FLOPs) scale. <https://openathena.ai/blog/quantile-balancing/>, April 2026. **Independent report of Quantile Balancing at 32B total / 5B active, 64 routed experts, 326B tokens. Blog post; no peer-reviewed source.**
- [52] OpenAI. Report on autonomous sandbox escape and external compromise during Exploit-Gym cyber-capability evaluation, July 2026. See also UK AI Safety Institute assessments of long-horizon cyber capability.
- [53] OpenRouter. Model routing: the Auto Router. <https://openrouter.ai/docs/features/model-routing>, 2026. **Deployed commercial system; vendor documentation, read directly.**
- [54] L. Beyer et al. PaliGemma: a versatile 3B VLM for transfer. *arXiv:2407.07726*, 2024.
- [55] S. Rajput et al. Recommender systems with generative retrieval (TIGER). *Advances in Neural Information Processing Systems*, 36, 2023.
- [56] Rijksinspectie Digitale Infrastructuur. Market surveillance under the AI Regulation. <https://www.rdi.nl/>, accessed 2026.

- [57] Dynamic model routing and cascading for efficient LLM inference: a survey. *arXiv:2603.04445*, 2026.
- [58] W. Sharon. Conditions for predictable social dynamics: conservation, decomposition, and control at criticality. Draft v0.5, position paper (in review), 2026.
- [59] Y. Sheng et al. S-LoRA: serving thousands of concurrent LoRA adapters. *Proceedings of MLSys*, 2024.
- [60] M. Shirokikh and S. Nikolenko. Sparse prefix caching for hybrid and recurrent LLM serving. *arXiv:2605.05219*, 2026.
- [61] I. Shumailov, Z. Shumaylov, Y. Zhao, N. Papernot, R. Anderson, and Y. Gal. AI models collapse when trained on recursively generated data. *Nature*, 631:755–759, 2024.
- [62] D. Silver and R. S. Sutton. Welcome to the era of experience. Google DeepMind, 2025.
- [63] H. A. Simon. The architecture of complexity. *Proceedings of the American Philosophical Society*, 106(6):467–482, 1962.
- [64] Multi-turn distributed inference with mixture of experts for 6G edge and cloud networks (StateFlow). *arXiv:2607.02522*, 2026.
- [65] J. Su. Blog post introducing Quantile Balancing for mixture-of-experts load balancing (in Chinese). <https://kexue.fm/archives/11619>, February 2026. **Personal blog; originating description of Quantile Balancing, predating the Kimi K3 release. The page returns HTTP 403 to automated retrieval, so we record the attribution as corroborated by [51] and by a public statement from P. Liang, rather than as read by us.**
- [66] J. Suarez. PufferLib: making reinforcement learning libraries and environments play nice. *arXiv:2406.12905*, 2024. Pokémon Red environment: PufferAI/pokegym.
- [67] S. Sukhbaatar, O. Golovneva, et al. Branch-Train-MiX: mixing expert LLMs into a mixture-of-experts LLM. *arXiv:2403.07816*; *COLM*, 2024.
- [68] W. Sun, X. Cheng, J. Fan, Y. Xu, X. Yu, S. He, J. Zhao, and K. Liu. Towards agentic self-learning LLMs in search environment. *arXiv:2510.14253*, 2025.
- [69] H. Sun, Y. Liu, Y. Wu, and L. Sun. Expert threshold routing for autoregressive language modeling with dynamic computation allocation and load balancing. *arXiv:2603.11535*, 2026.
- [70] Moonshot AI (Kimi Team). Kimi Linear: an expressive, efficient attention architecture. *arXiv:2510.26692*, 2025. Introduces Kimi Delta Attention.
- [71] Tracer AI. Echo: adaptive coordination of open-weight models. <https://tracerm1.ai/>, 2026. **Vendor-reported; no independent evaluation located.**
- [72] P. Villalobos, A. Ho, J. Sevilla, T. Besiroglu, L. Heim, and M. Hobbhahn. Will we run out of data? Limits of LLM scaling based on human-generated data. *arXiv:2211.04325v2*, 2024; *Proceedings of ICML* (Position track), 2024.
- [73] vLLM Project. A preview of production-scale Kimi K3 support on vLLM. <https://vllm.ai/blog/2026-07-22-kimi-k3-preview>, July 2026.
- [74] A. Wang et al. HMoE: heterogeneous mixture of experts for language modeling. *arXiv:2408.10681*, 2024.

- [75] L. Wang et al. Auxiliary-loss-free load balancing strategy for mixture-of-experts. *arXiv:2408.15664*, 2024.
- [76] Y. Yue et al. Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model? *Advances in Neural Information Processing Systems*, 38, 2025. Best Paper Runner-up.
- [77] H. Zenil. On the limits of self-improving in large language models: the singularity is not near without symbolic model synthesis. *arXiv:2601.05280*, 2026.
- [78] H. Zhang et al. MiLoRA: efficient mixture of low-rank adaptation for large language models fine-tuning. *Findings of EMNLP*, 2024.
- [79] A. L. Zhang, T. L. Griffiths, K. R. Narasimhan, and O. Press. VideoGameBench: can vision-language models complete popular video games? *arXiv:2505.18134*, 2025.
- [80] X. Zhu, C. Zhang, Y. Chi, T. Stafford, N. Collier, and A. Vlachos. Demystifying multi-agent debate: the role of confidence and diversity. *arXiv:2601.19921*, 2026. **Proves the confidence-weighted submartingale (Thm. 1); its diversity result changes the initial distribution, not the update dynamics. No author named Choi is on this paper; it is not choi2025debatevote.**
- [81] A. Zhao et al. Absolute Zero: reinforced self-play reasoning with zero data. *arXiv:2505.03335*; *Advances in Neural Information Processing Systems*, 38, 2025.
- [82] Zhipu AI. GLM-5.2 open-weight release, June 2026. <https://huggingface.co/zai-org/GLM-5.2> (MIT licence).